



**KOÇ
UNIVERSITY**

Database Management Systems

Relational Model

M. Emre Gürsoy

Assistant Professor
Department of Computer Engineering

www.memregursoy.com



Introduction

- The **relational model** was first proposed by Dr. E. F. Codd from IBM Research in 1970:
 - "A relational model for large shared data banks", Communications of the ACM, June 1970
 - Revolutionary paper in the field of DBMS
 - Dr. Codd won the ACM Turing Award
- The **relational model** is based on the concept of a **relation**
 - Informally, a **relation** looks like a **table** with many **rows**
 - "Header" at the top (**schema** of the relation)
 - Rows are data entries (**tuples** in the relation)



A "Relation"

- The **schema** of a relation specifies how it is defined
 - Name of relation, name of attributes, attribute domains
 - Contents (tuples) are not included in the schema
- We typically write: **R(A1, A2, ..., An)**
 - Relation name **R**, attributes **A1, A2, ..., An**
 - Example: **STUDENT(Name, Ssn, Home_phone, Address, Office_phone, Age, Gpa)**

Relation Name		Attributes						
STUDENT		Name	Ssn	Home_phone	Address	Office_phone	Age	Gpa
Tuples	→	Benjamin Bayer	305-61-2435	373-1616	2918 Bluebonnet Lane	NULL	19	3.21
	→	Chung-cha Kim	381-62-1245	375-4409	125 Kirby Road	NULL	18	2.89
	→	Dick Davidson	422-11-2320	NULL	3452 Elgin Road	749-1253	25	3.53
	→	Rohan Panchal	489-22-1100	376-9821	265 Lark Lane	749-6492	28	3.93
	→	Barbara Benson	533-69-1238	839-8461	7384 Fontana Lane	NULL	19	3.25



Formal Definitions

- Formally, given $R(A_1, A_2, \dots, A_n)$
 - R is the **name** of the relation
 - A_1, A_2, \dots, A_n are the **attributes** of the relation
 - $R(A_1, A_2, \dots, A_n)$ is the **schema** of the relation
 - Each attribute has a "**domain**" of allowed values:
 $\text{dom}(A_j)$
- $r(R)$: a specific **state** of relation R – this is a *set of tuples* (rows) that currently exist in the relation
 - $r(R) = \{t_1, t_2, \dots\}$ where each t_i is a tuple
 - $t_i = \langle v_1, v_2, \dots, v_n \rangle$ where each $v_j \in \text{dom}(A_j)$
 - $r(R) \subseteq \text{dom}(A_1) \times \text{dom}(A_2) \times \dots \times \text{dom}(A_n)$



Formal Definitions - Example

- Let $R(A1, A2)$ be a relation schema:
 - Let $\text{dom}(A1) = \{0,1\}$
 - Let $\text{dom}(A2) = \{a,b,c\}$
- Then: $\text{dom}(A1) \times \text{dom}(A2)$ is all possible combinations:
 $\{ \langle 0,a \rangle, \langle 0,b \rangle, \langle 0,c \rangle, \langle 1,a \rangle, \langle 1,b \rangle, \langle 1,c \rangle \}$
- The relation state $r(R) \subseteq \text{dom}(A1) \times \text{dom}(A2)$
- For example: $r(R)$ could be $\{ \langle 0,a \rangle, \langle 0,b \rangle, \langle 1,c \rangle \}$
 - This is one possible state (or “population”) r of relation R
 - It has three tuples: $\langle 0,a \rangle, \langle 0,b \rangle, \langle 1,c \rangle$



Tuples

- We use notation $t[A_j]$ to mean the value tuple t has for attribute A_j .
- A special **NULL** value is used to represent values that are unknown/unavailable in the database.
 - Below, for some tuples, $t[\text{Office_phone}] = \text{NULL}$.
- Tuples in a relation are **not ordered** (why?)
 - I can shuffle the table arbitrarily, $r(R)$ stays the same
- All tuples in a relation are **unique** (why?)

STUDENT

Name	Ssn	Home_phone	Address	Office_phone	Age	Gpa
Dick Davidson	422-11-2320	NULL	3452 Elgin Road	749-1253	25	3.53
Barbara Benson	533-69-1238	839-8461	7384 Fontana Lane	NULL	19	3.25
Rohan Panchal	489-22-1100	376-9821	265 Lark Lane	749-6492	28	3.93
Chung-cha Kim	381-62-1245	375-4409	125 Kirby Road	NULL	18	2.89
Benjamin Bayer	305-61-2435	373-1616	2918 Bluebonnet Lane	NULL	19	3.21



Formal vs Informal

<u>Informal Terms</u>		<u>Formal Terms</u>
Table		Relation
Column Header		Attribute
All possible values for that column		Domain of Attribute
Row		Tuple
Table Definition		Schema of a Relation
Populated Table		State of the Relation



Relational Database

- Collection of several relations (schema+state)
- The schema of Company DB:

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------

DEPT_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
----------------	------------------

PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
-------	----------------	-----------	------

WORKS_ON

<u>Essn</u>	<u>Pno</u>	Hours
-------------	------------	-------

DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
-------------	-----------------------	-----	-------	--------------

Figure 5.5
Schema diagram for
the COMPANY
relational database
schema.



Key + Superkey

- **Superkey** **SK** of relation R is a set of attributes such that:
 - Two tuples cannot have the same value for all **SK** attributes, i.e., for distinct **t1**, **t2**: $t1[SK] \neq t2[SK]$
 - This condition must hold for **all** valid states of R
- **Key** **K** of relation R:
 - A "minimal" superkey
 - If you remove any attribute from **K**, uniqueness no longer holds

STUDENT

Name	Ssn	Home_phone	Address	Office_phone	Age	Gpa
Dick Davidson	422-11-2320	NULL	3452 Elgin Road	749-1253	25	3.53
Barbara Benson	533-69-1238	839-8461	7384 Fontana Lane	NULL	19	3.25
Rohan Panchal	489-22-1100	376-9821	265 Lark Lane	749-6492	28	3.93
Chung-cha Kim	381-62-1245	375-4409	125 Kirby Road	NULL	18	2.89
Benjamin Bayer	305-61-2435	373-1616	2918 Bluebonnet Lane	NULL	19	3.21



Key + Superkey

- True/False?
 - A key is always a superkey.
 - A superkey is always a key.
 - Any set of attributes that includes a key is a superkey.
- Can there be multiple keys?
 - CAR(State, Reg#, SerialNo, Make, Model, Year)
 - CAR has two keys:
 - Key1 = {State, Reg#}
 - Key2 = {SerialNo}
 - Both are also superkeys
 - {SerialNo, Make} is a superkey but not a key




Primary Key

- If a relation has several candidate keys (eg: CAR), one of them is chosen (by the administrator) as the **primary key**.
- We **underline the primary key** in the schema.
 - CAR(State, Reg#, SerialNo, Make, Model, Year)
 - STUDENT(Name, SSN, Home_phone, Address, Office_phone, Age, GPA)
- The primary key is used to uniquely identify each tuple in a relation (it becomes the "identity" of the tuple).
- Recommendation (but not a rule) for choosing the primary key: **Choose the smallest of the candidate keys.**
 - Key1 = {State, Reg#} has 2 attributes
 - Key2 = {SerialNo} has 1 attribute



Foreign Key

- In a DB with many relations, attributes in some relations may reference attributes in other relations.
- Tuples in the referencing relation (eg: ENROLLMENT) have attributes **FK** called **foreign keys** that reference the primary key attributes **PK** of the referenced relation (eg: STUDENTS).
 - **sid** is a primary key in STUDENTS; it is a foreign key in ENROLLMENT





Examples of Foreign Keys

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------

DEPT_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
----------------	------------------

PROJECT

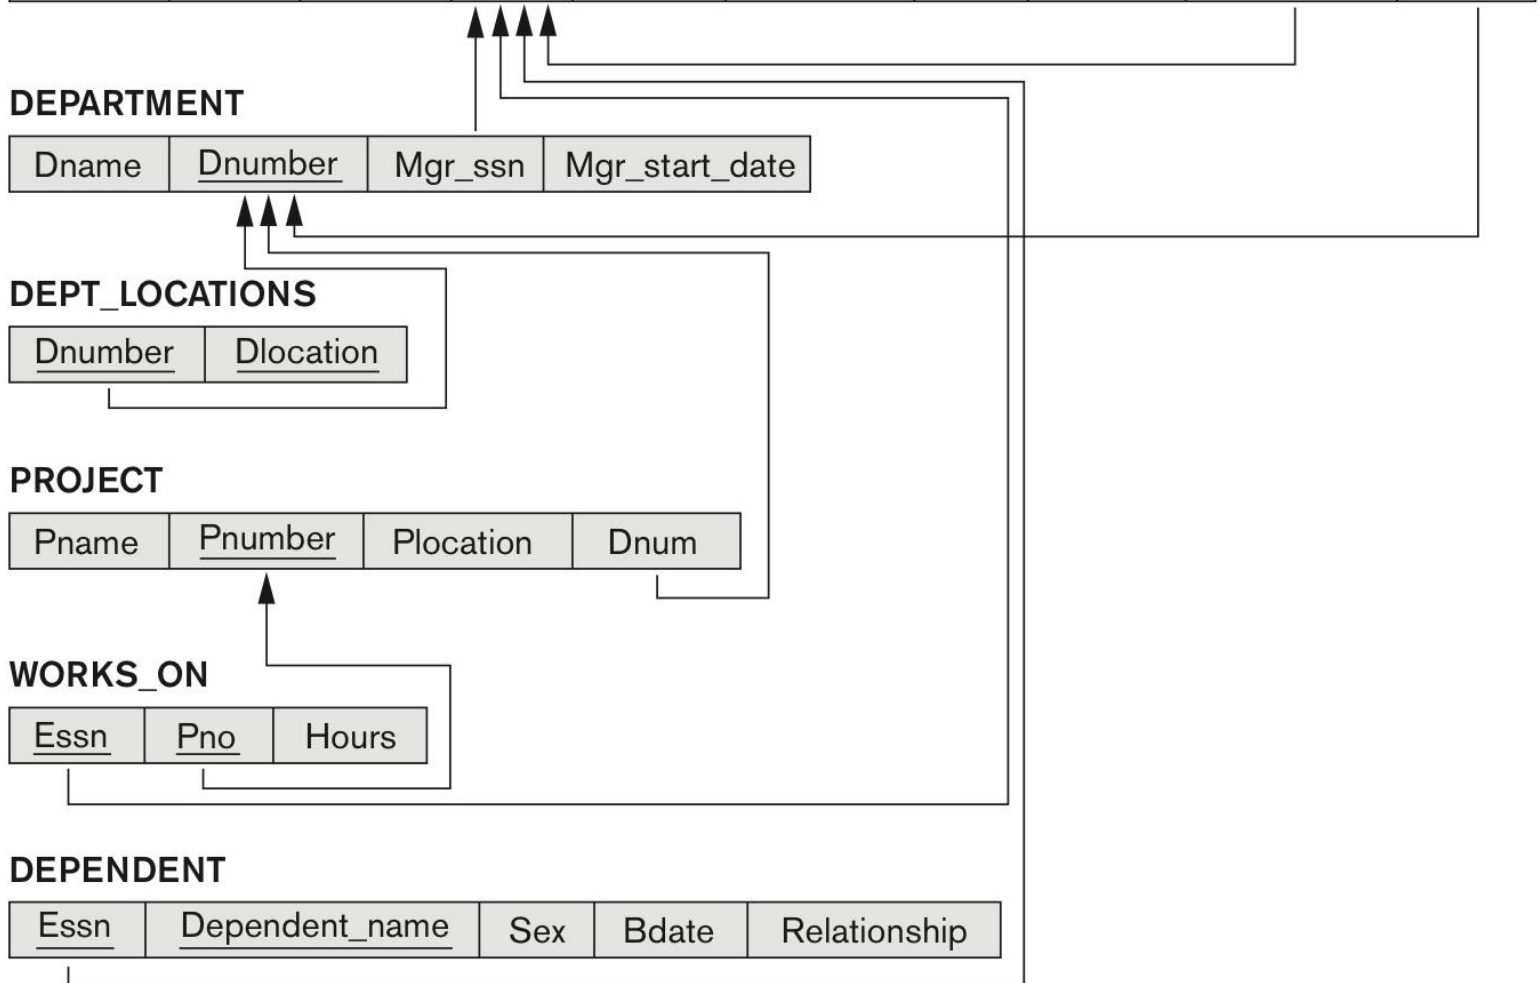
Pname	<u>Pnumber</u>	Plocation	Dnum
-------	----------------	-----------	------

WORKS_ON

<u>Essn</u>	<u>Pno</u>	Hours
-------------	------------	-------

DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
-------------	-----------------------	-----	-------	--------------





Company DB Example

Figure 5.6

One possible database state for the COMPANY relational database schema.

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

DEPARTMENT

Dname	<u>Dnumber</u>	<u>Mgr_ssn</u>	<u>Mgr_start_date</u>
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

DEPT_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
1	Houston
4	Stafford
5	Bellaire
5	Sugarland
5	Houston

WORKS_ON

<u>Essn</u>	<u>Pno</u>	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0
999887777	30	30.0
999887777	10	10.0
987987987	10	35.0
987987987	30	5.0
987654321	30	20.0
987654321	20	15.0
888665555	20	NULL

PROJECT

<u>Pname</u>	<u>Pnumber</u>	<u>Plocation</u>	<u>Dnum</u>
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
333445555	Alice	F	1986-04-05	Daughter
333445555	Theodore	M	1983-10-25	Son
333445555	Joy	F	1958-05-03	Spouse
987654321	Abner	M	1942-02-28	Spouse
123456789	Michael	M	1988-01-04	Son
123456789	Alice	F	1988-12-30	Daughter
123456789	Elizabeth	F	1967-05-05	Spouse



Creating Relations

```
CREATE TABLE Students  
(sid: CHAR(20),  
name: CHAR(20),  
login: CHAR(10),  
age: INTEGER,  
gpa: REAL,  
PRIMARY KEY (sid))
```

<u>sid</u>	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Shero	shero@cs	18	3.2

```
CREATE TABLE Enrollment  
(sid: CHAR(20),  
cid: CHAR(30),  
grade: CHAR(2),  
PRIMARY KEY (sid, cid),  
FOREIGN KEY (sid)  
REFERENCES Students)
```

<u>sid</u>	<u>cid</u>	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B



Integrity Constraints

- 4 types of integrity constraints in the relational model:
 - **Domain constraint:** Every value of a tuple must be from the domain of its corresponding attribute (or NULL, if allowed).
 - **Key constraint:** Every tuple must be unique in terms of its key attributes.
 - **Entity integrity constraint:** Tuples cannot have NULL values for their primary key attributes.
 - **Referential integrity constraint:** There should be no dangling references between relations.
 - More detail on the next slides!



Referential Integrity

- What should be done if someone wants to...
 - Insert a tuple into Enrollment with a non-existent **sid**?
 - Reject it!
 - Delete a tuple from Students?
 - Also delete all Enrollment tuples that refer to that student tuple?
 - Disallow deletion of this tuple?
 - Update sid of a student in Students?
 - Disallow?
 - Modify that sid in Enrollment?

```
CREATE TABLE Enrollment
(sid: CHAR(20),
cid: CHAR(30),
grade: CHAR(2),
PRIMARY KEY (sid, cid),
FOREIGN KEY (sid)
REFERENCES Students)
```

<u>sid</u>	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Shero	shero@cs	18	3.2

<u>sid</u>	<u>cid</u>	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B



Referential Integrity

- All of these are examples of how referential integrity (RI) could potentially be violated.
- Example ways in which RI violations can be handled:
 - **NO ACTION:** reject deletion/update
 - **CASCADE:** also delete/update all tuples that refer to the deleted/updated tuple
 - **SET DEFAULT:** set foreign key value to a pre-determined, special DEFAULT value

```
CREATE TABLE Enrollment
(sid: CHAR(20),
cid: CHAR(30),
grade: CHAR(2),
PRIMARY KEY (sid, cid),
FOREIGN KEY (sid)
REFERENCES Students
ON DELETE CASCADE
ON UPDATE SET DEFAULT)
```



Exercise

- Consider a relational database with many relations
- Consider 3 main kinds of operations that arise in practice:
 - INSERT a tuple into a relation
 - DELETE a tuple from a relation
 - MODIFY/UPDATE a tuple in a relation
- For each of these operations, which integrity constraints can they possibly violate? Examples?