



**KOÇ
UNIVERSITY**

Database Management Systems

Concurrency Control

M. Emre Gürsoy

Assistant Professor
Department of Computer Engineering

www.memregursoy.com

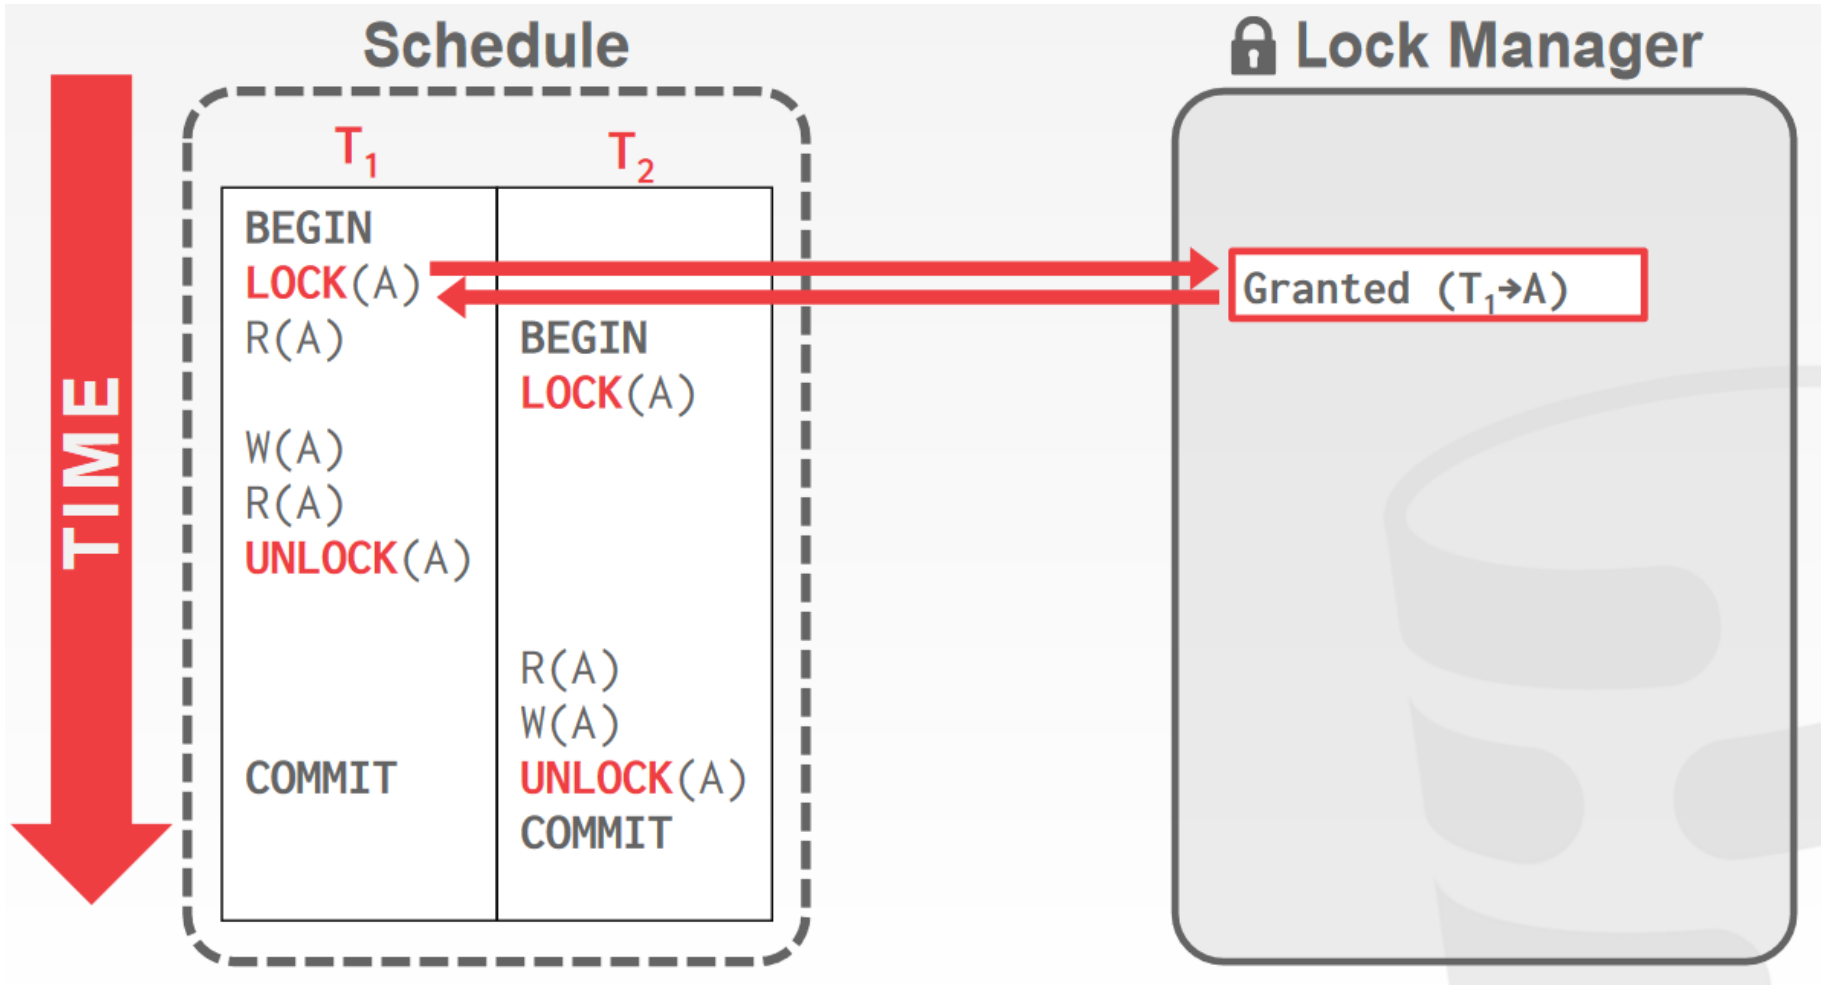


Introduction

- We have been talking about:
 - "Good" schedules versus "bad" schedules (e.g., conflict serializability)
 - How to check if a schedule is "good" or "bad"
- But our previous methods have two shortcomings:
 - They require us to know all transactions and all actions ahead of time (all **R**s and **W**s)
 - They don't tell us how to **create** a "good" schedule
- Now we'll learn about **locking** for **concurrency control**

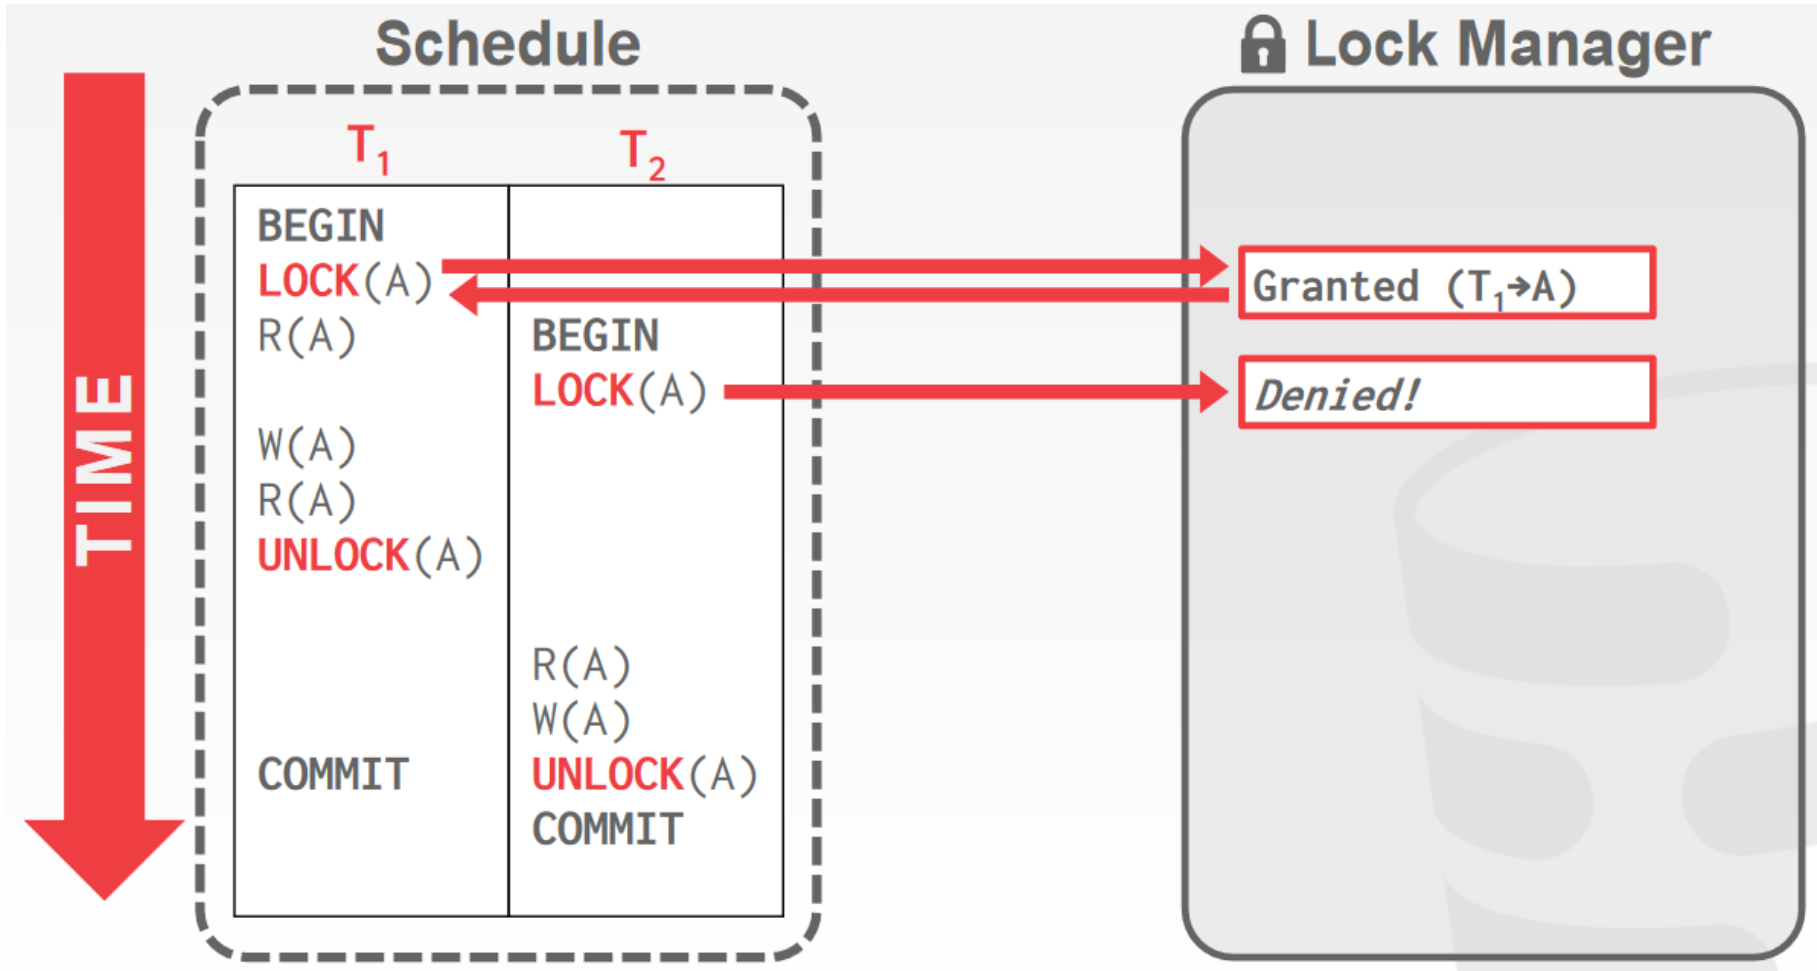


Intuition



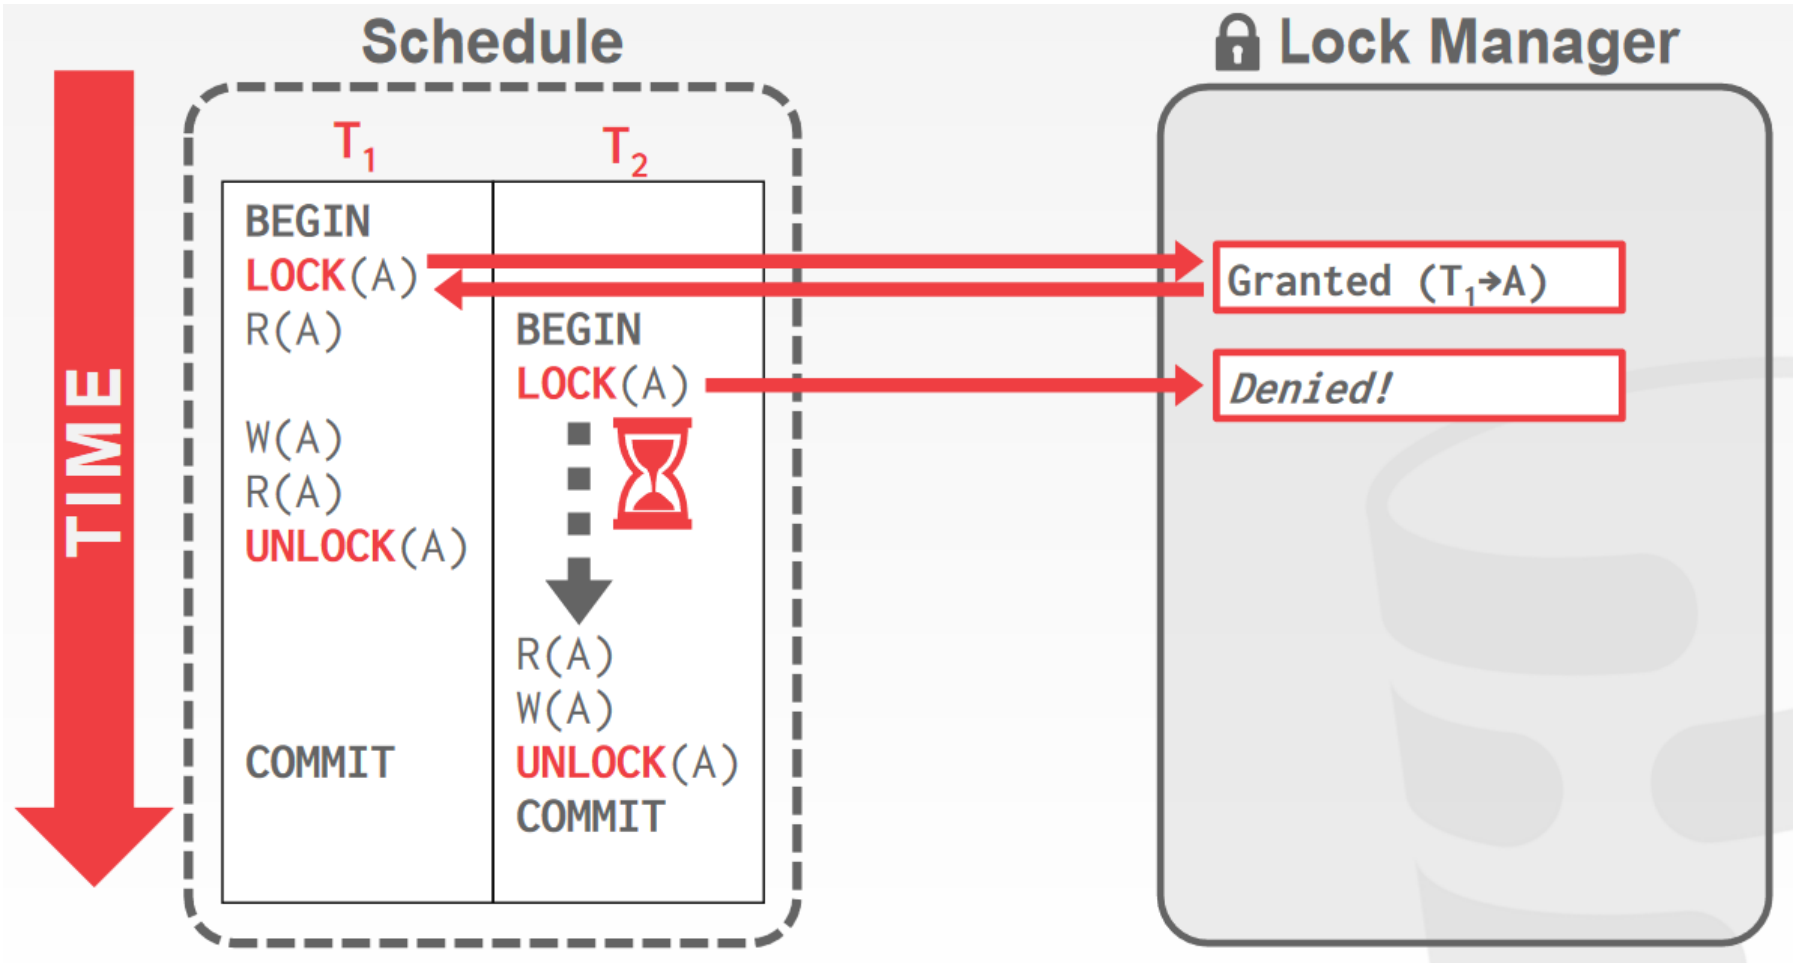


Intuition



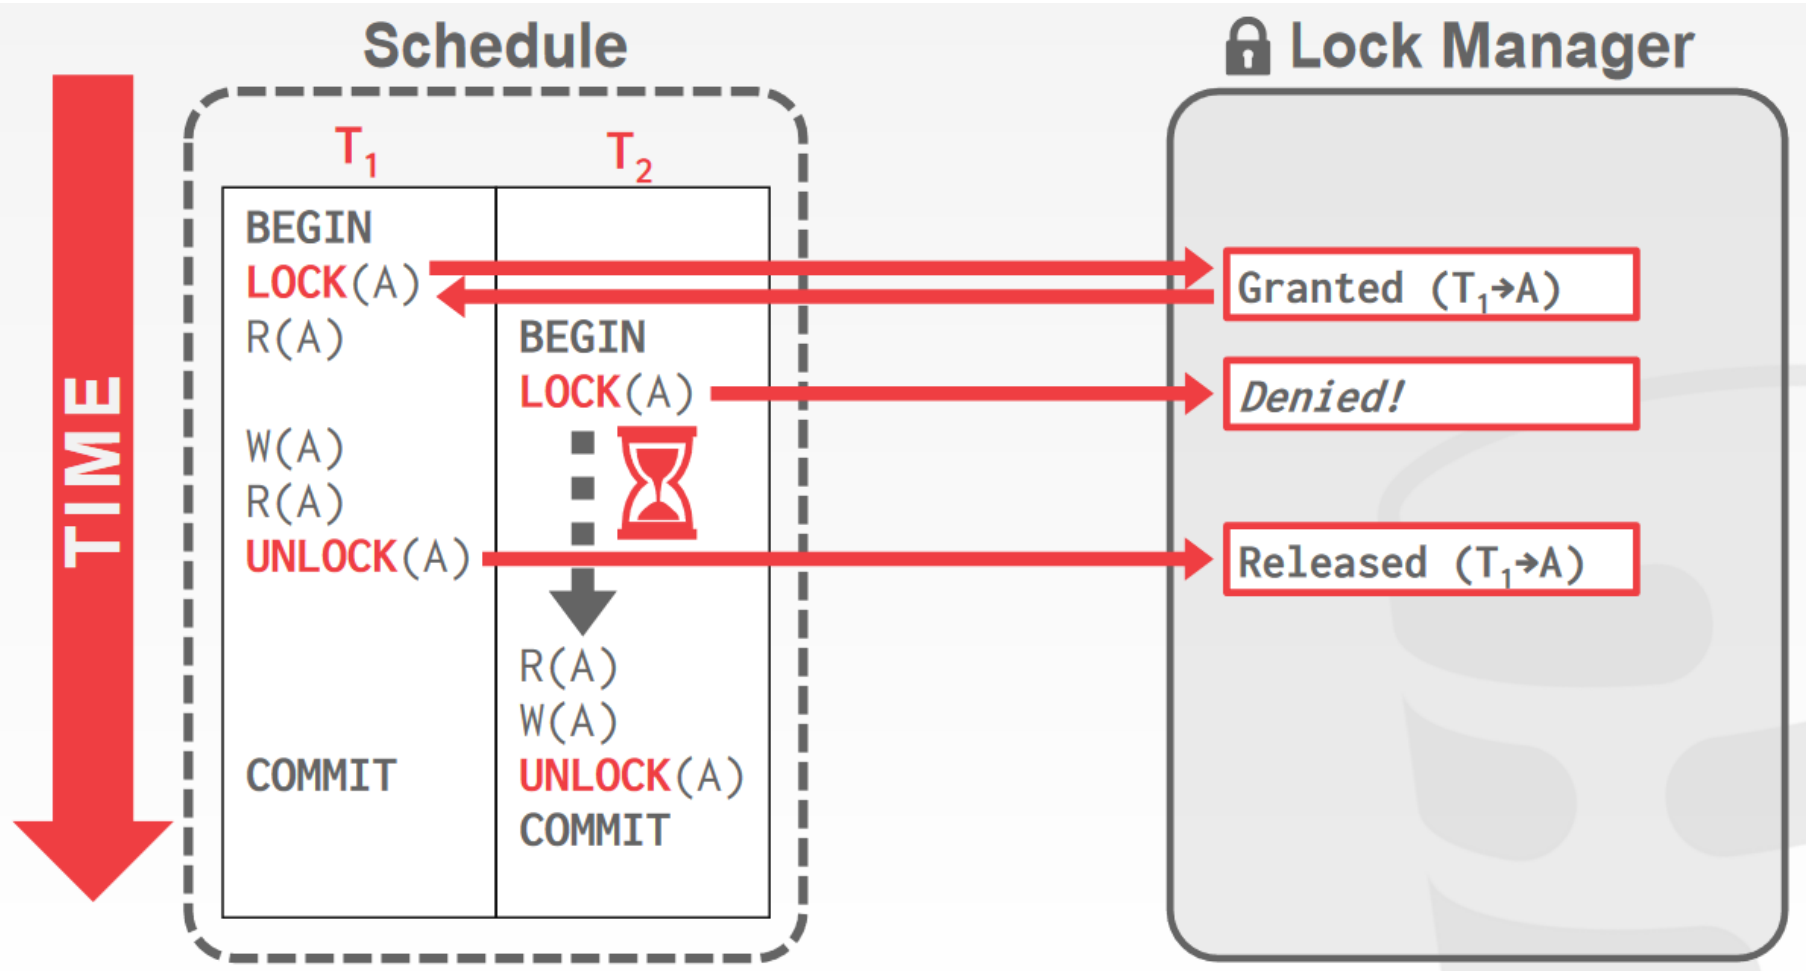


Intuition



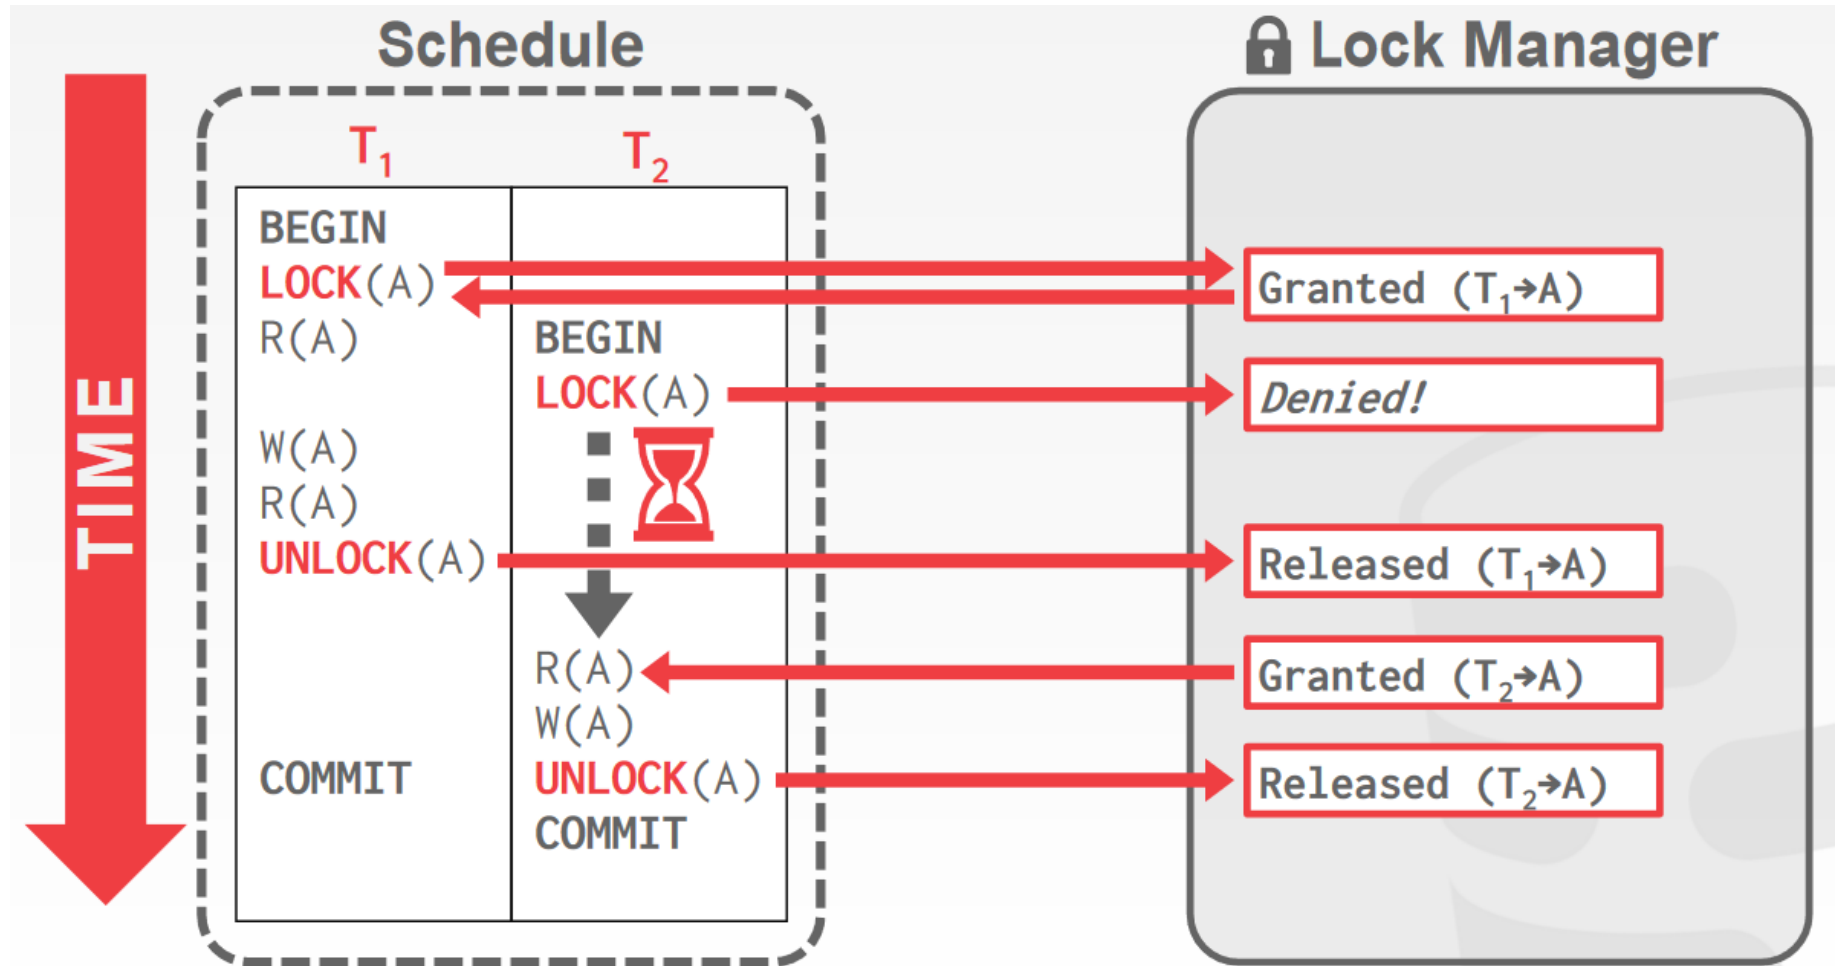


Intuition





Intuition





Lock Types

- Two types of locks:
 - Shared locks (S-lock):** for **reads**
 - Exclusive locks (X-lock):** for **writes**
- Why not use shared locks for writes?
- Why not use exclusive locks for reads?

Compatibility Matrix		
	Shared	Exclusive
Shared	✓	X
Exclusive	X	X

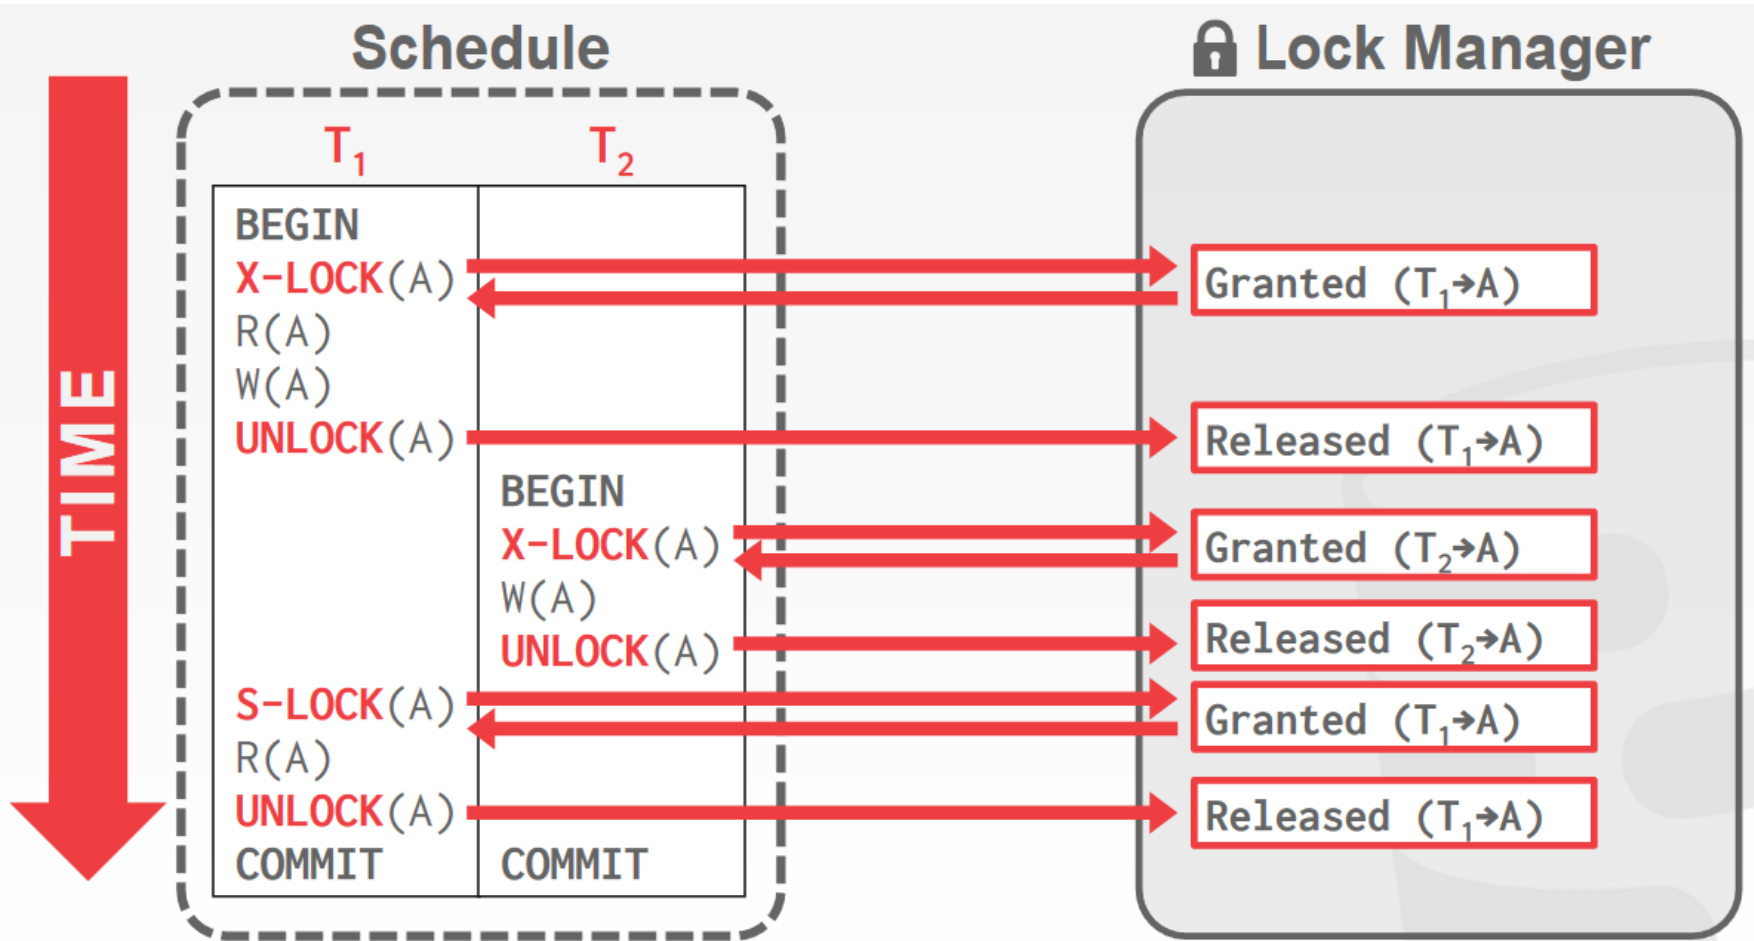


Executing with Locks

- Execution using locks:
 - Transactions **request** locks.
 - Lock manager **grants** or **denies** requests.
 - Once transactions are done with the DB object, they **release** their locks (**unlock**).
 - **Lock manager** keeps track of:
 - which transactions hold which locks on which objects
 - which transactions are waiting to acquire locks



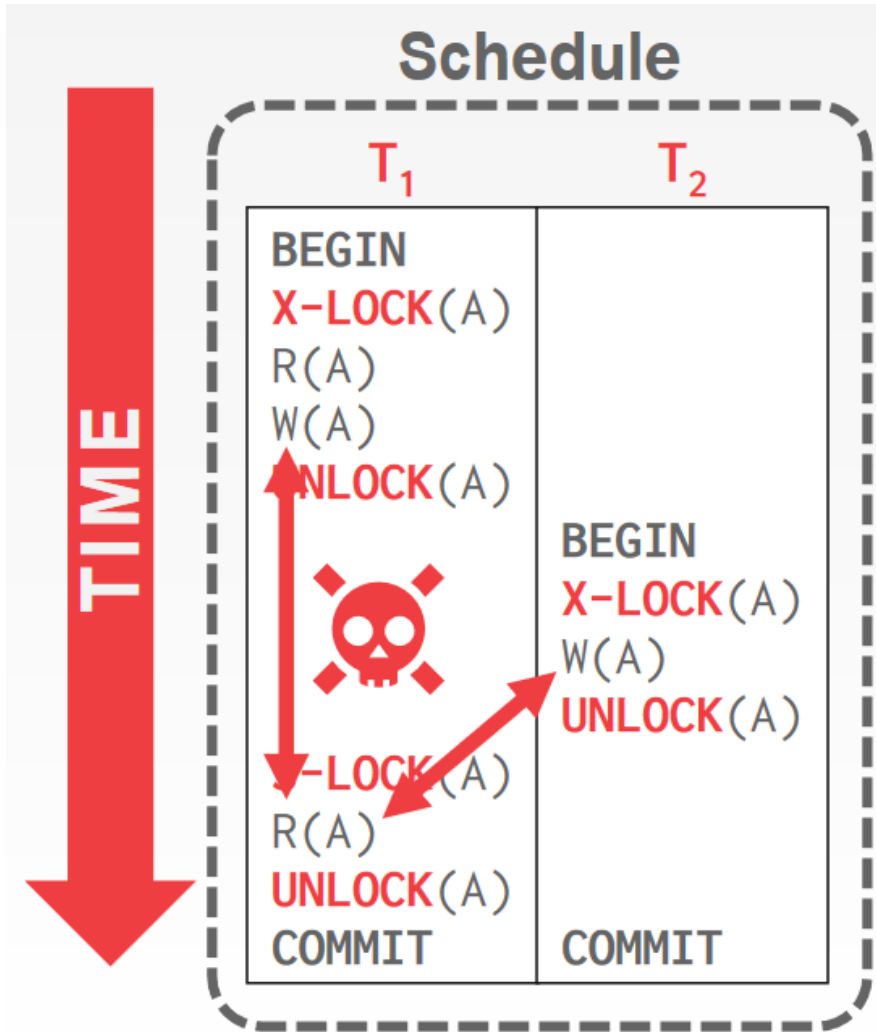
Executing with Locks



Do you see a problem here?



Executing with Locks



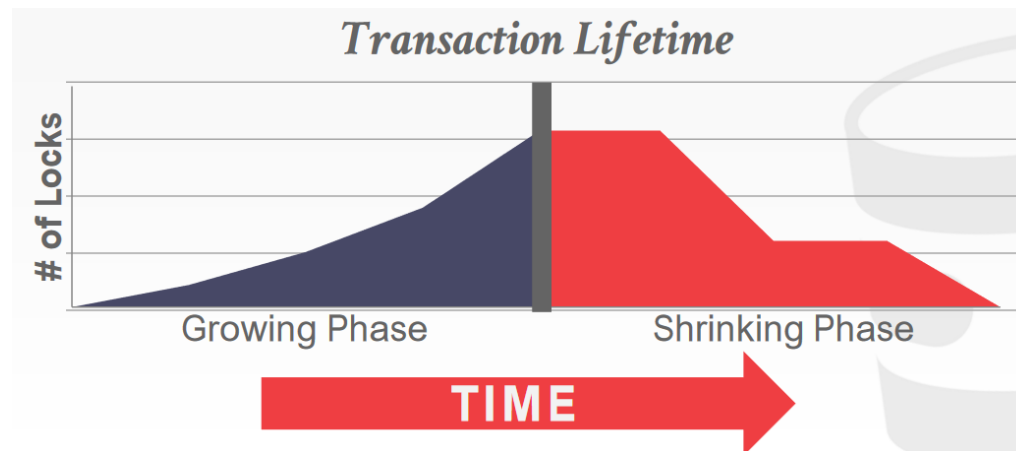
T_1 is expecting to read what it wrote but ends up reading what T_2 wrote.

What if T_2 decides to abort later?



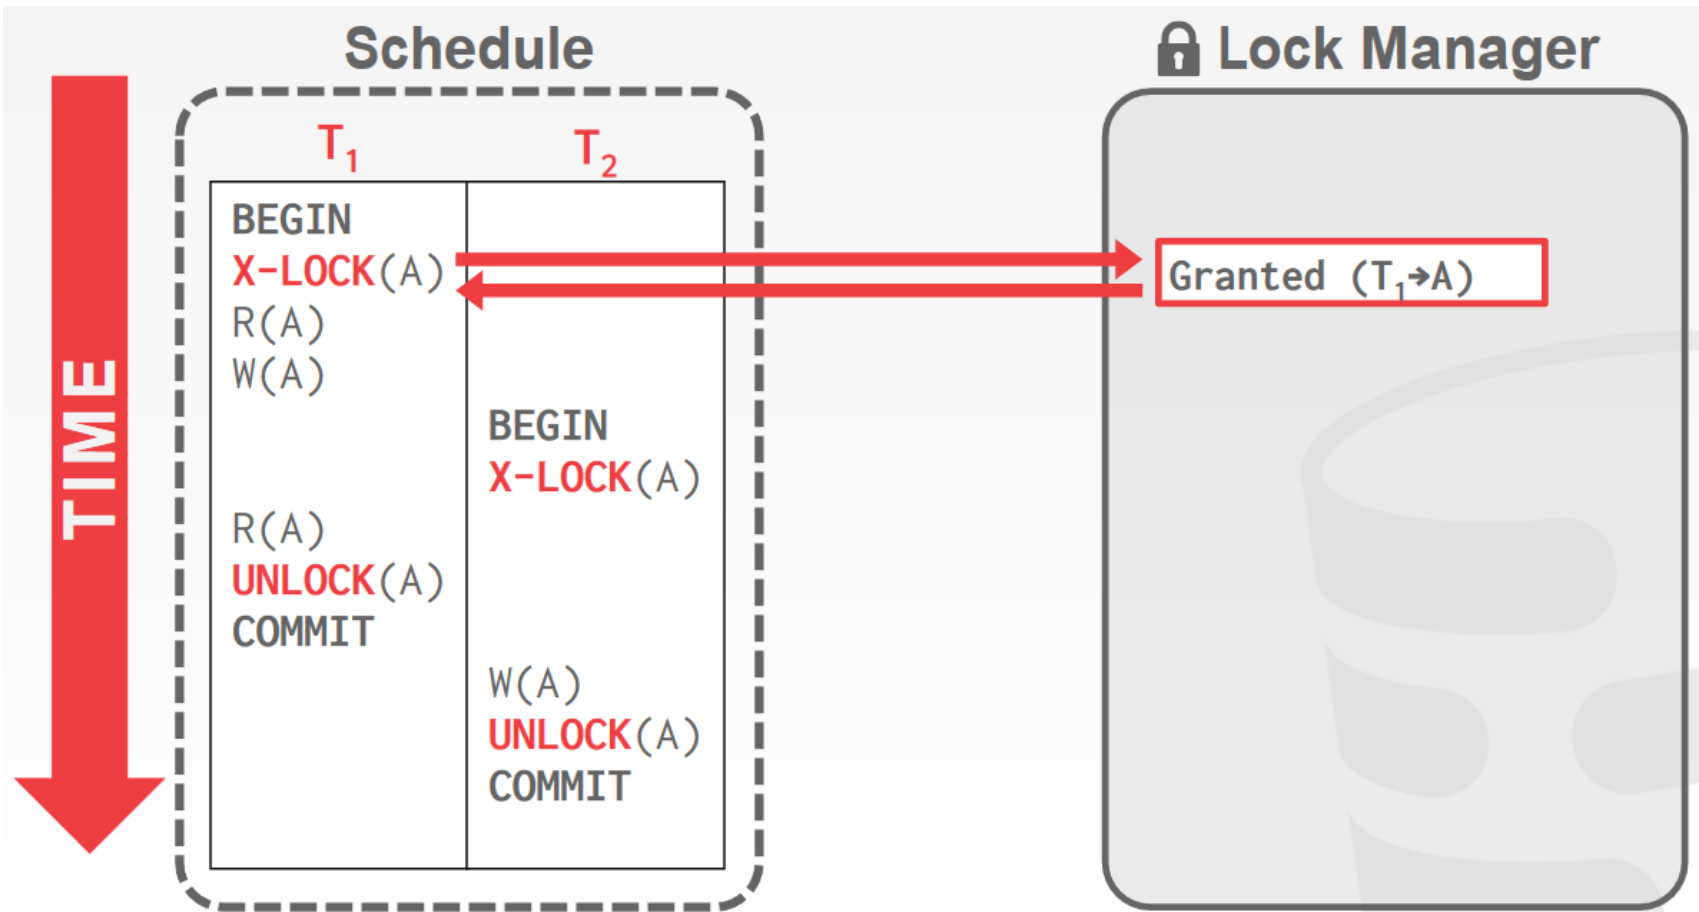
Two-Phase Locking

- **Two-phase locking (2PL)** is a commonly used concurrency control protocol.
 - Phase #1: **Growing**
 - Each transaction requests locks from DBMS lock manager.
 - Lock manager grants or denies lock requests.
 - Phase #2: **Shrinking**
 - The transaction is allowed to release previously acquired locks, but it is not allowed to acquire new locks.





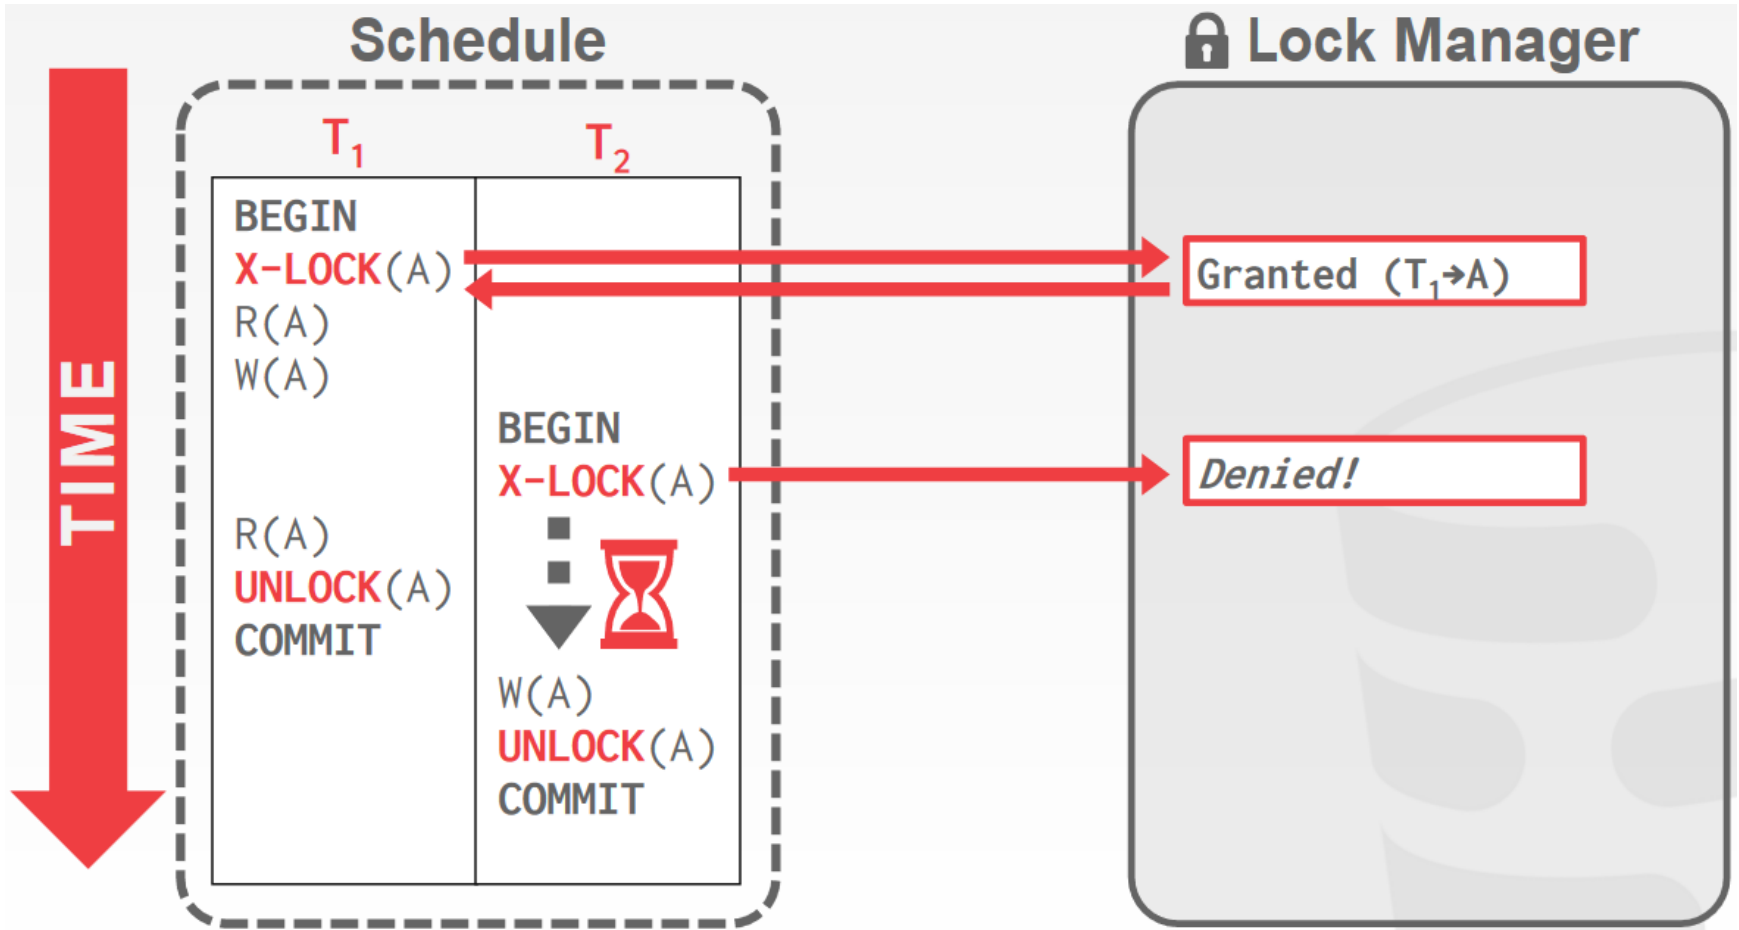
Executing with 2PL



What is the main difference with the previous case?

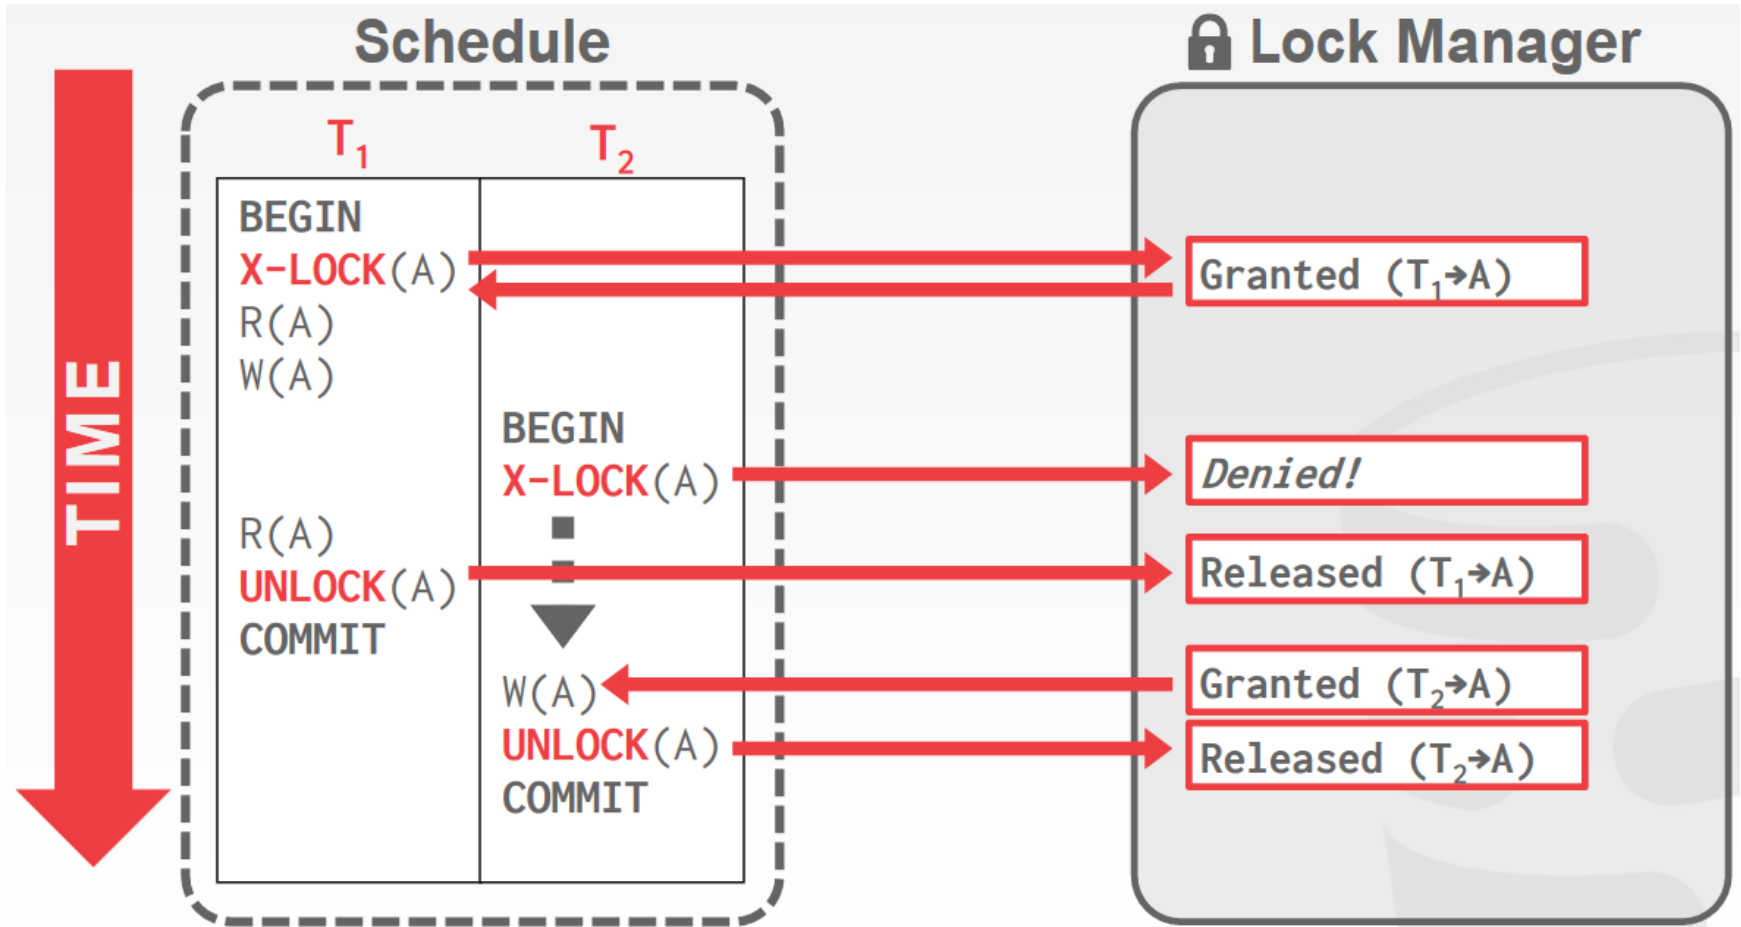


Executing with 2PL





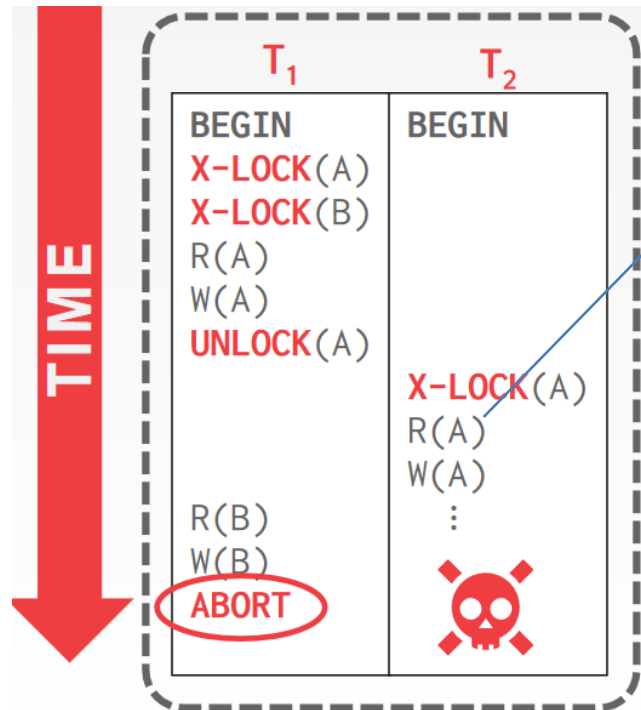
Executing with 2PL





Cascading Aborts

- 2PL is sufficient to guarantee conflict serializability.
- But 2PL is subject to the **cascading aborts** problem.
 - When a transaction aborts, it causes other transactions to also have to abort.




But T_2 had read what T_1 wrote and assumed that it was true!

So, when T_1 aborts we must also abort T_2 .



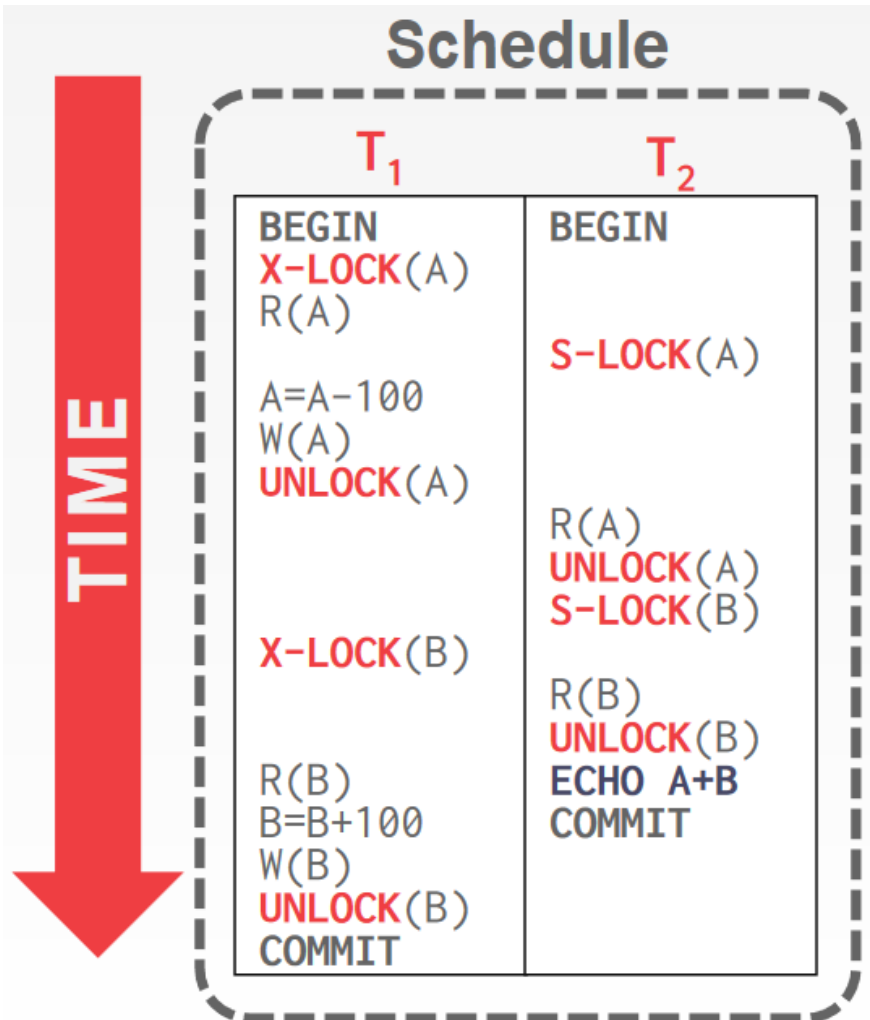
Strict 2PL

- All locks held by a transaction are released when the transaction completes.
 - No **unlocking** in the middle of a transaction.
 - If I have an **X-lock** on **A**, no one else can get an **S-lock** or **X-lock** on **A** before I commit or abort.
 - Hence, no one can read/write on **A** before I am finished.
 - Hence, no cascading aborts.





Example (Non-2PL)

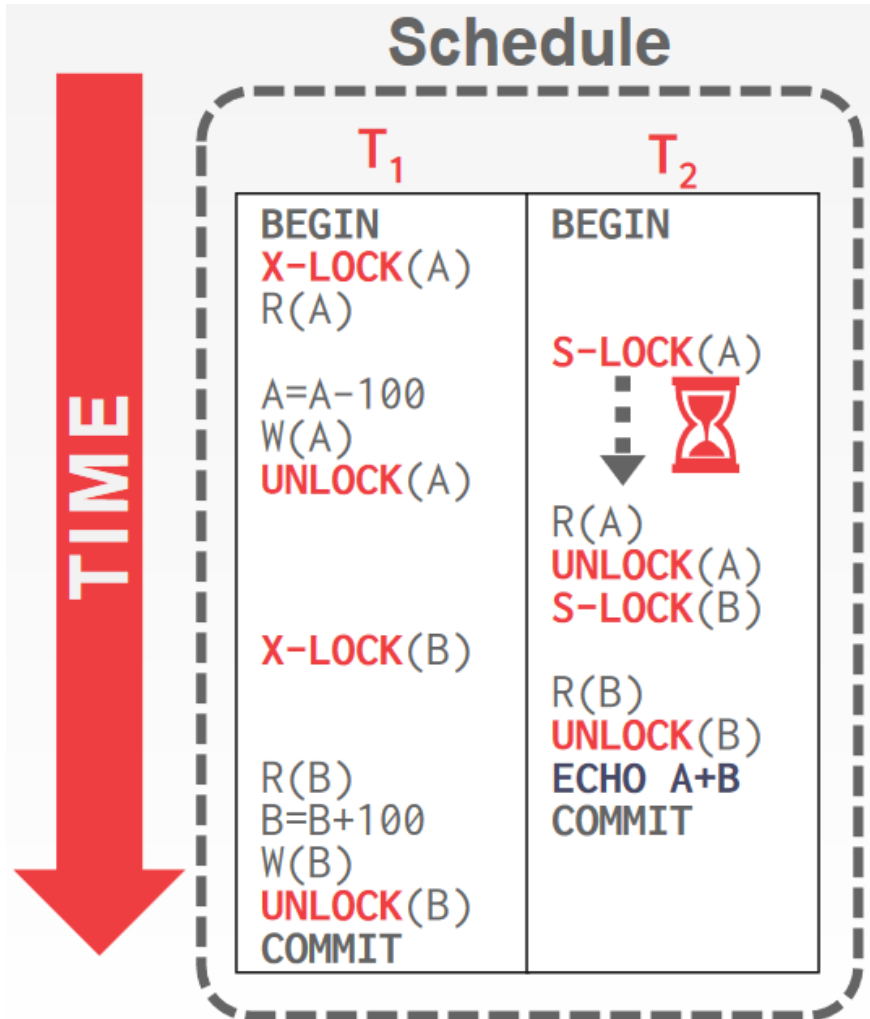


Assume $A=1000$, $B=1000$ initially.

What are the two transactions intending to do?

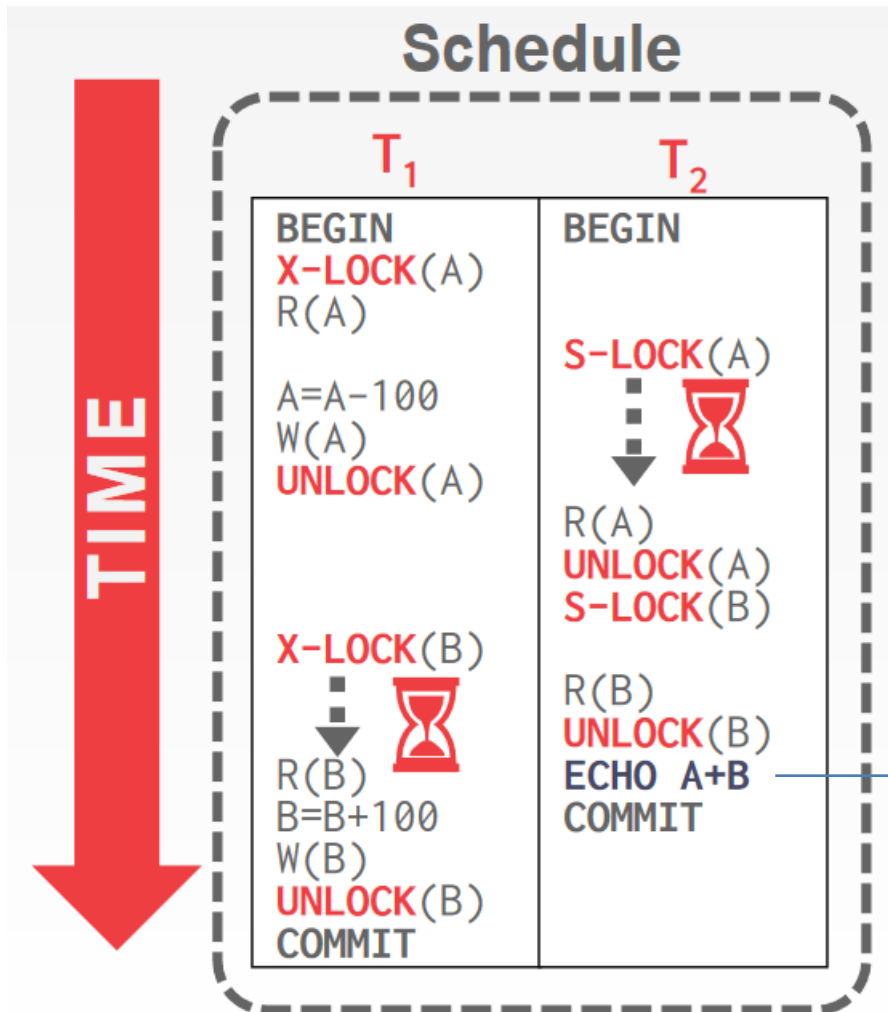


Example (Non-2PL)





Example (Non-2PL)



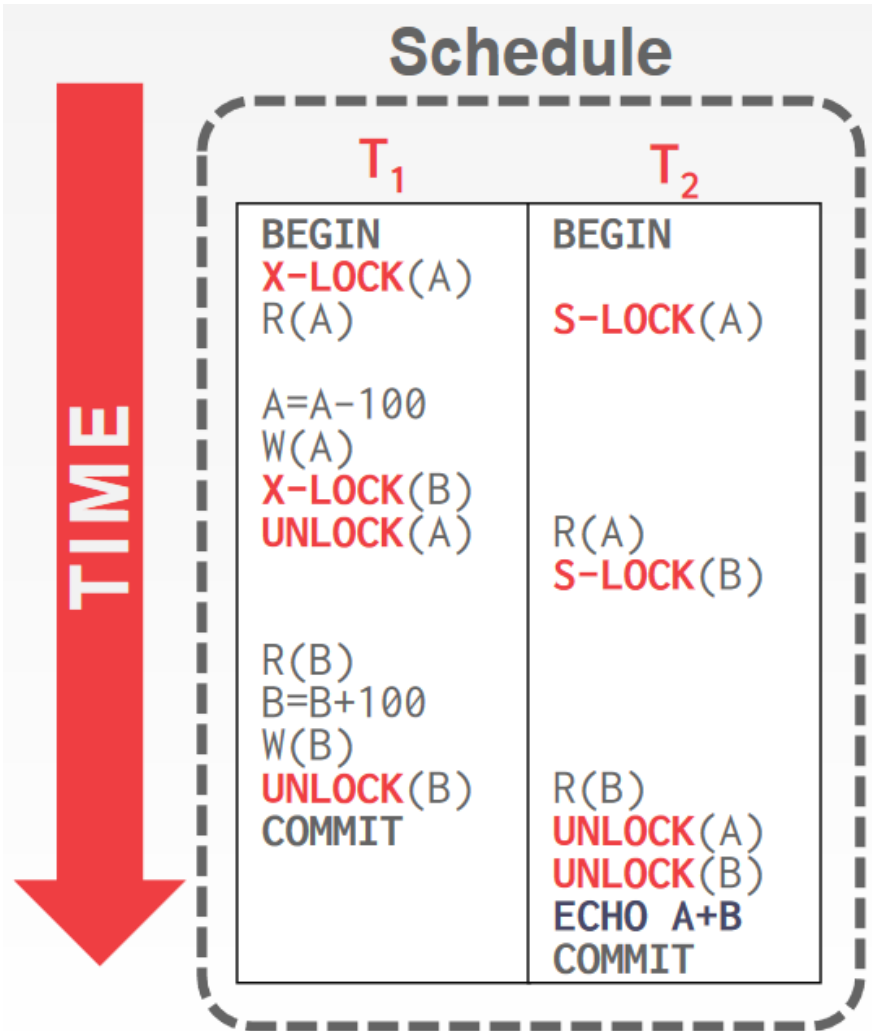
Assume $A=1000$, $B=1000$ initially.

What gets printed out?

What are the final values of A and B ?



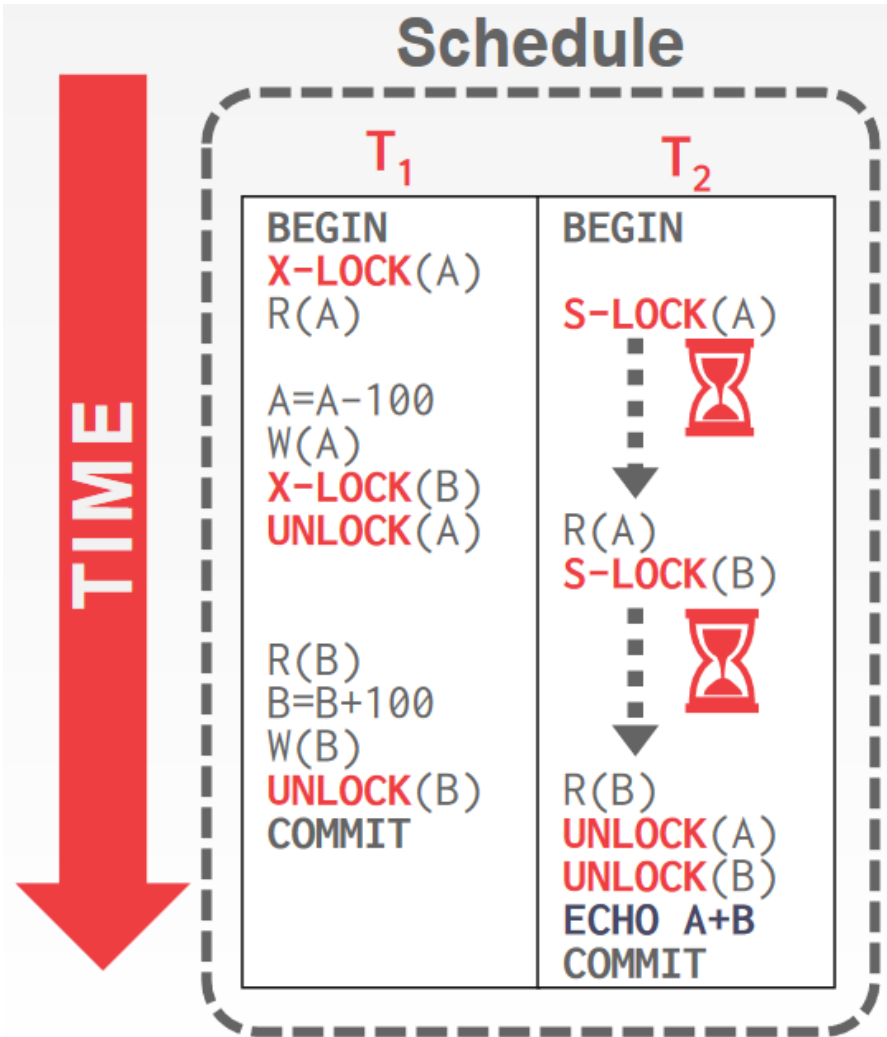
Example (2PL)



Assume $A=1000$, $B=1000$ initially.

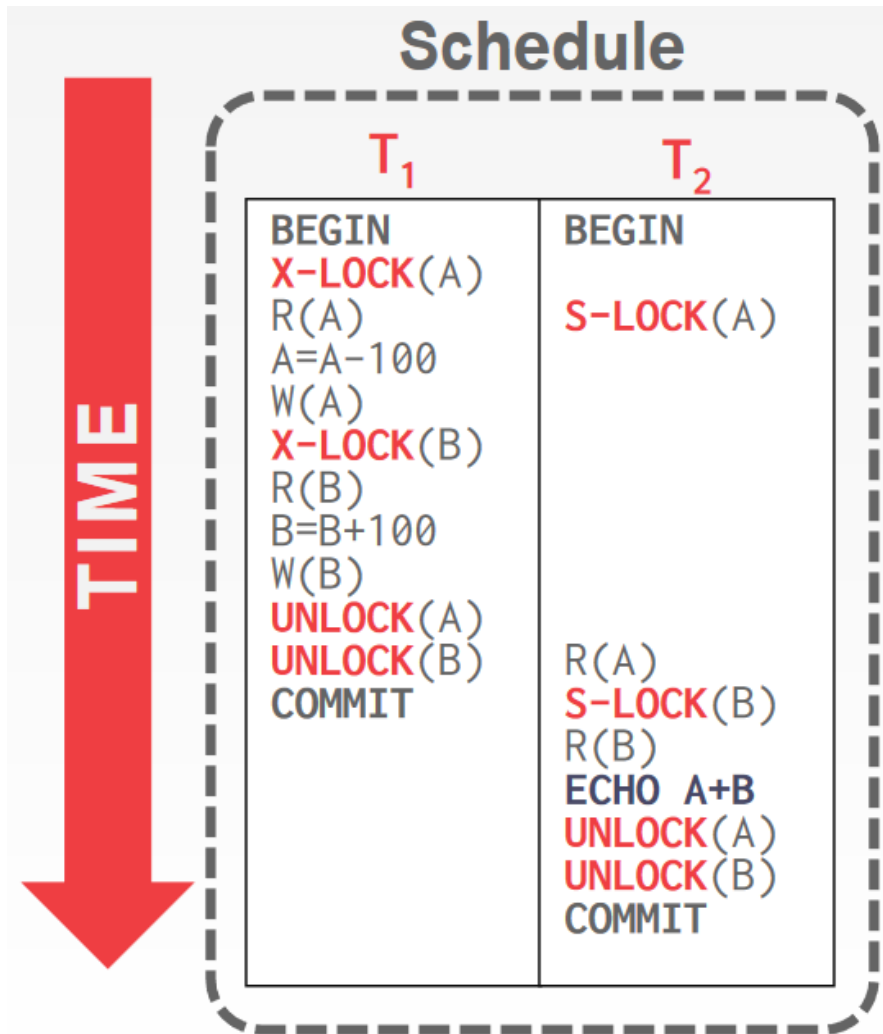


Example (2PL)





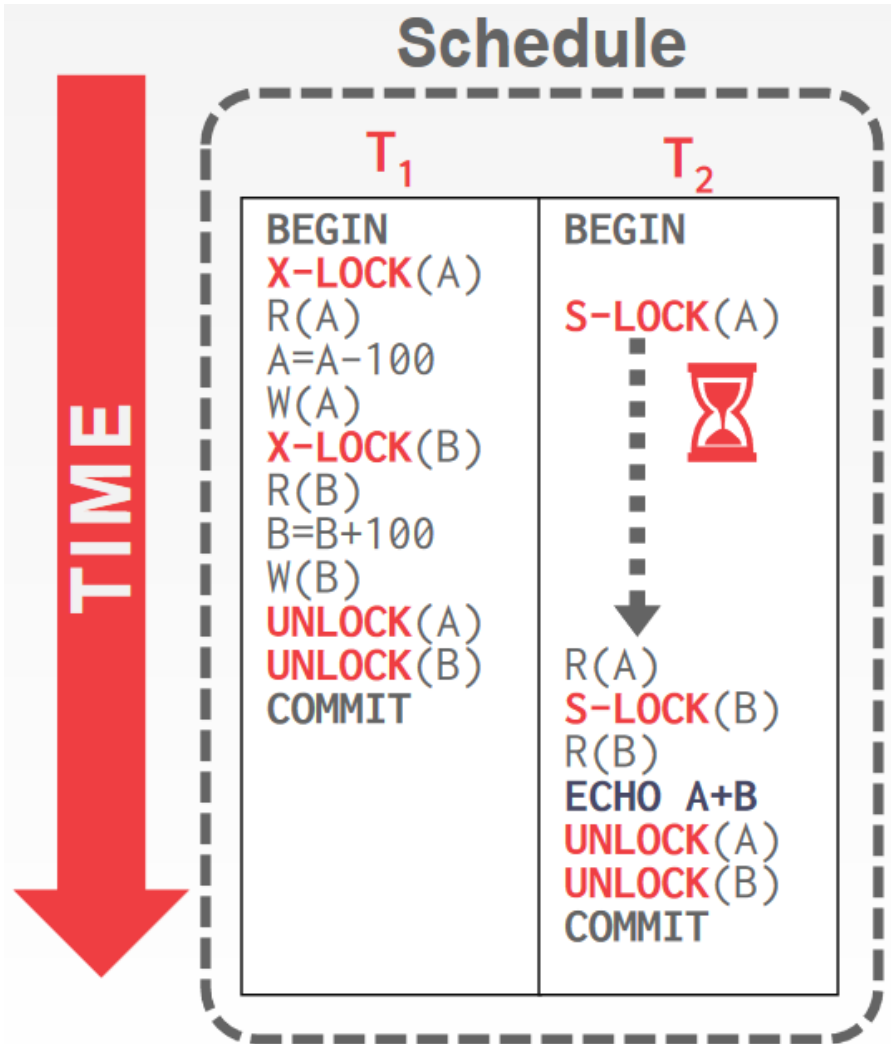
Example (Strict 2PL)



Assume $A=1000$, $B=1000$ initially.



Example (Strict 2PL)



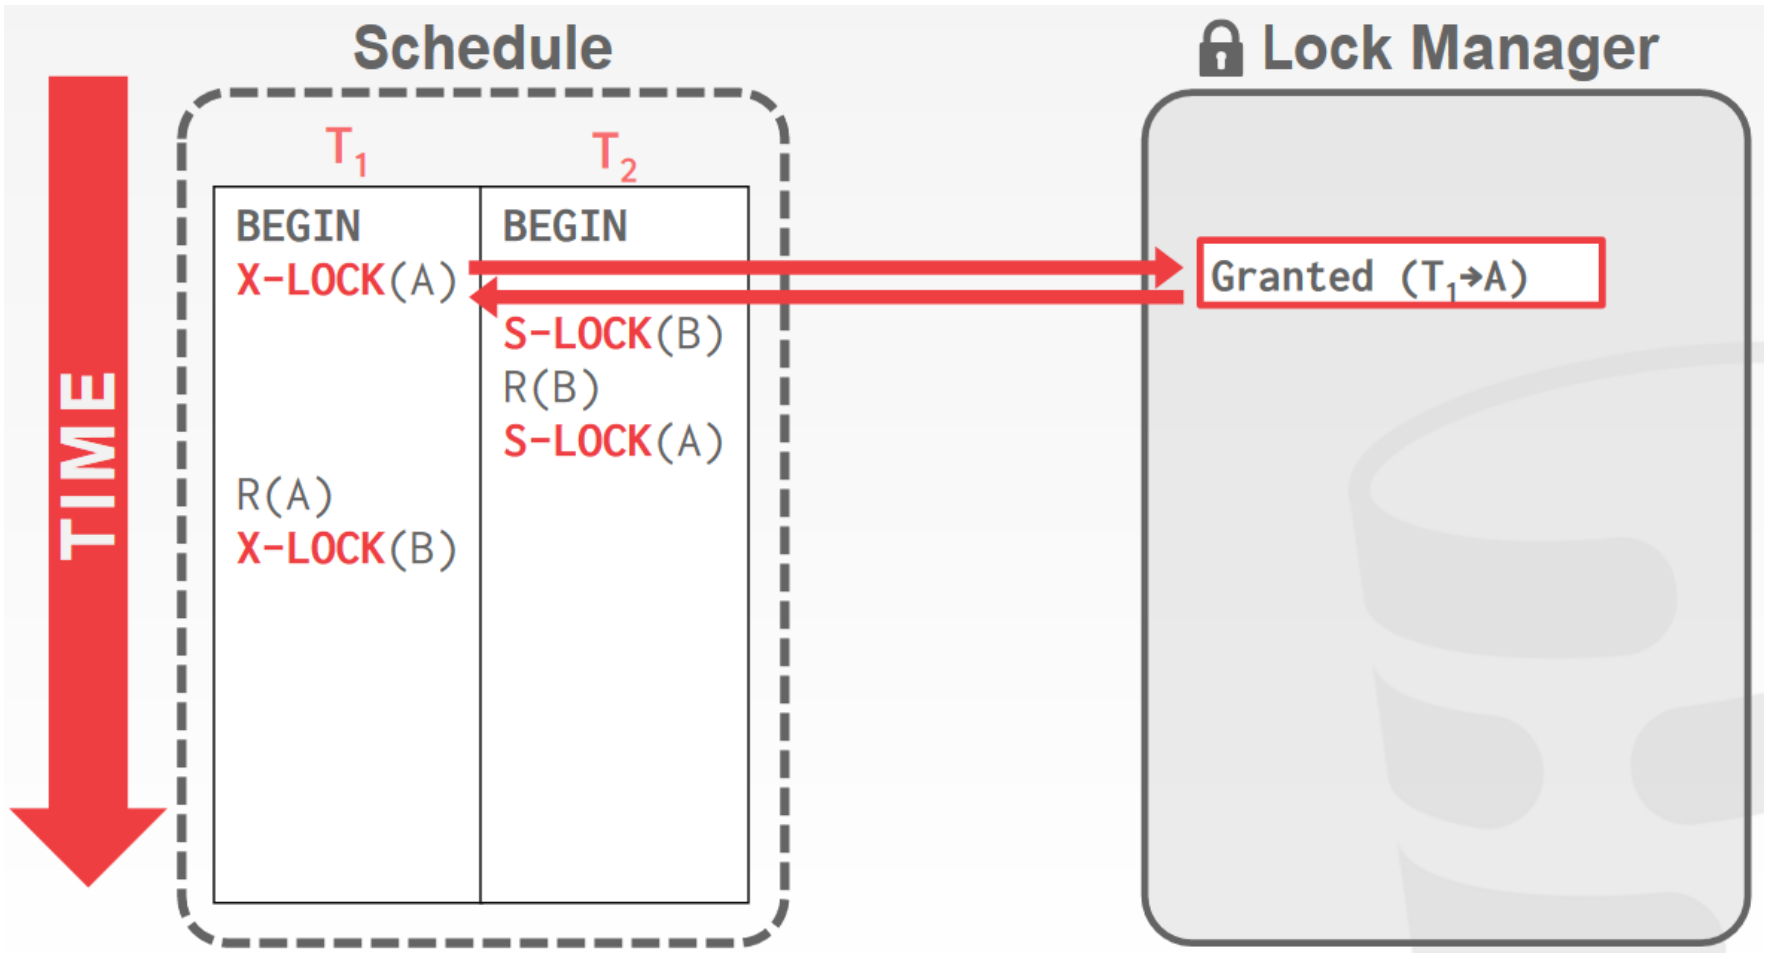


Deadlocks

- 2PL can lead to deadlocks.
- A **deadlock** is a cycle of transactions waiting for locks to be released by each other.
 - T_1 locks **A**, needs access to **B** in order to continue.
 - T_2 locks **B**, needs access to **A** in order to continue.
- Two issues related to deadlocks:
 - Deadlock detection
 - Deadlock prevention

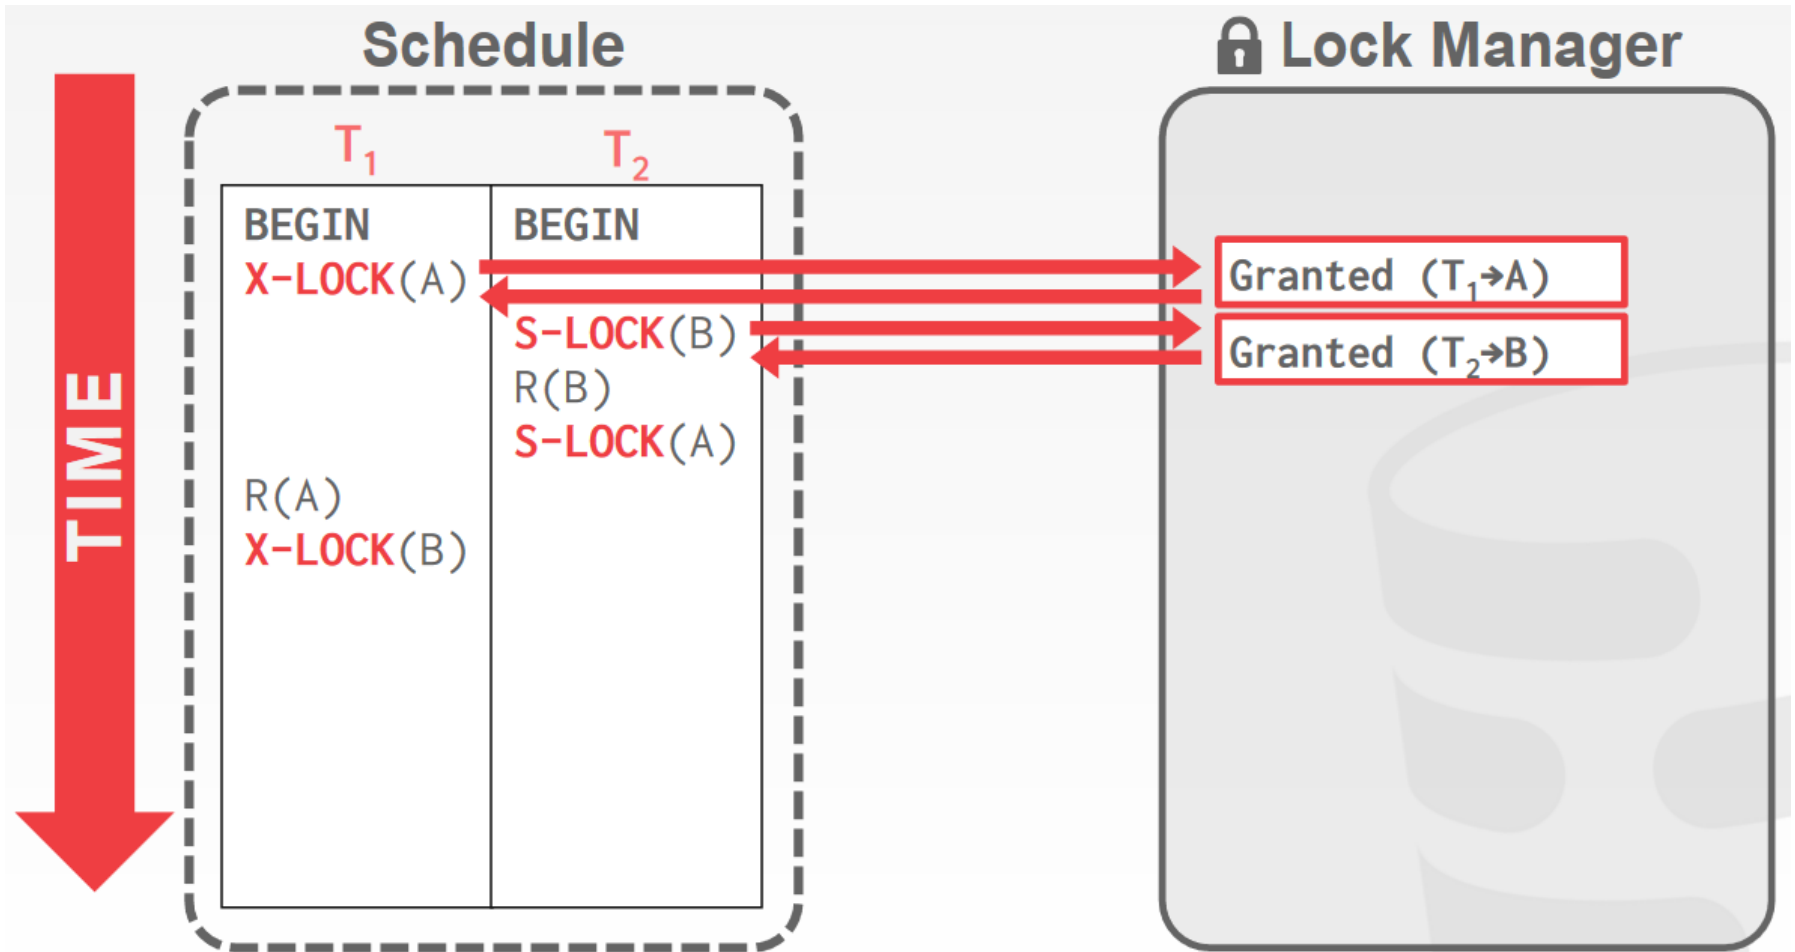


Deadlocks



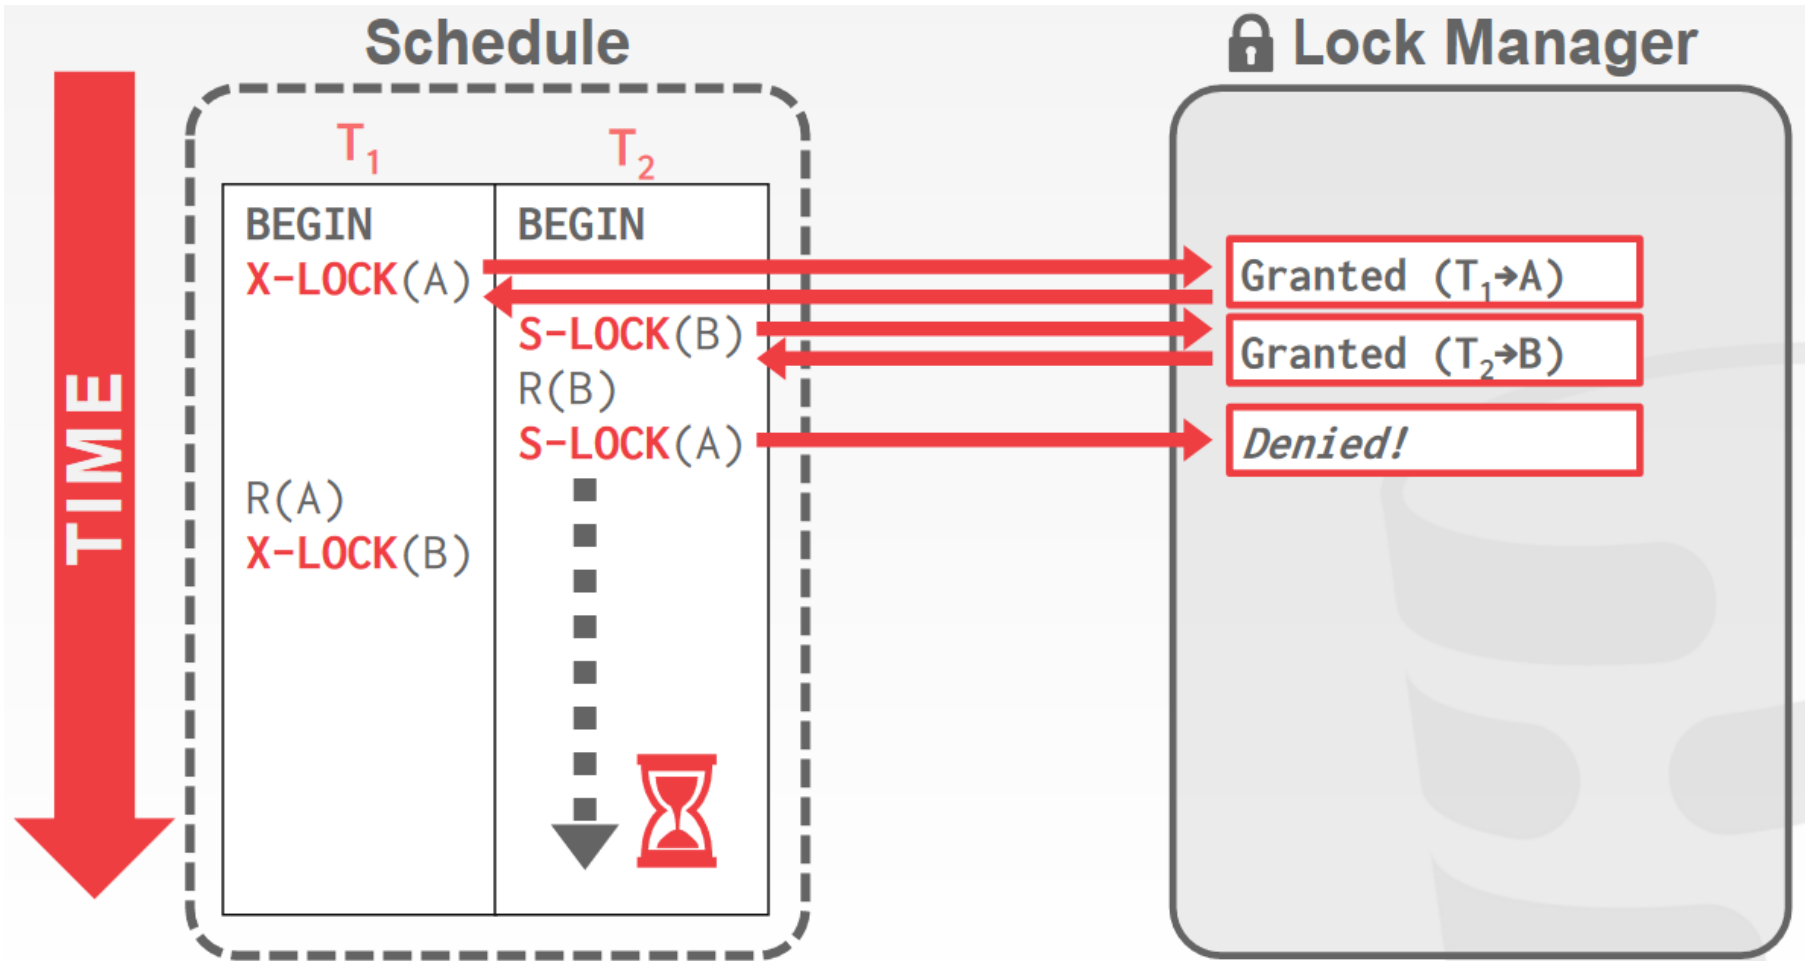


Deadlocks



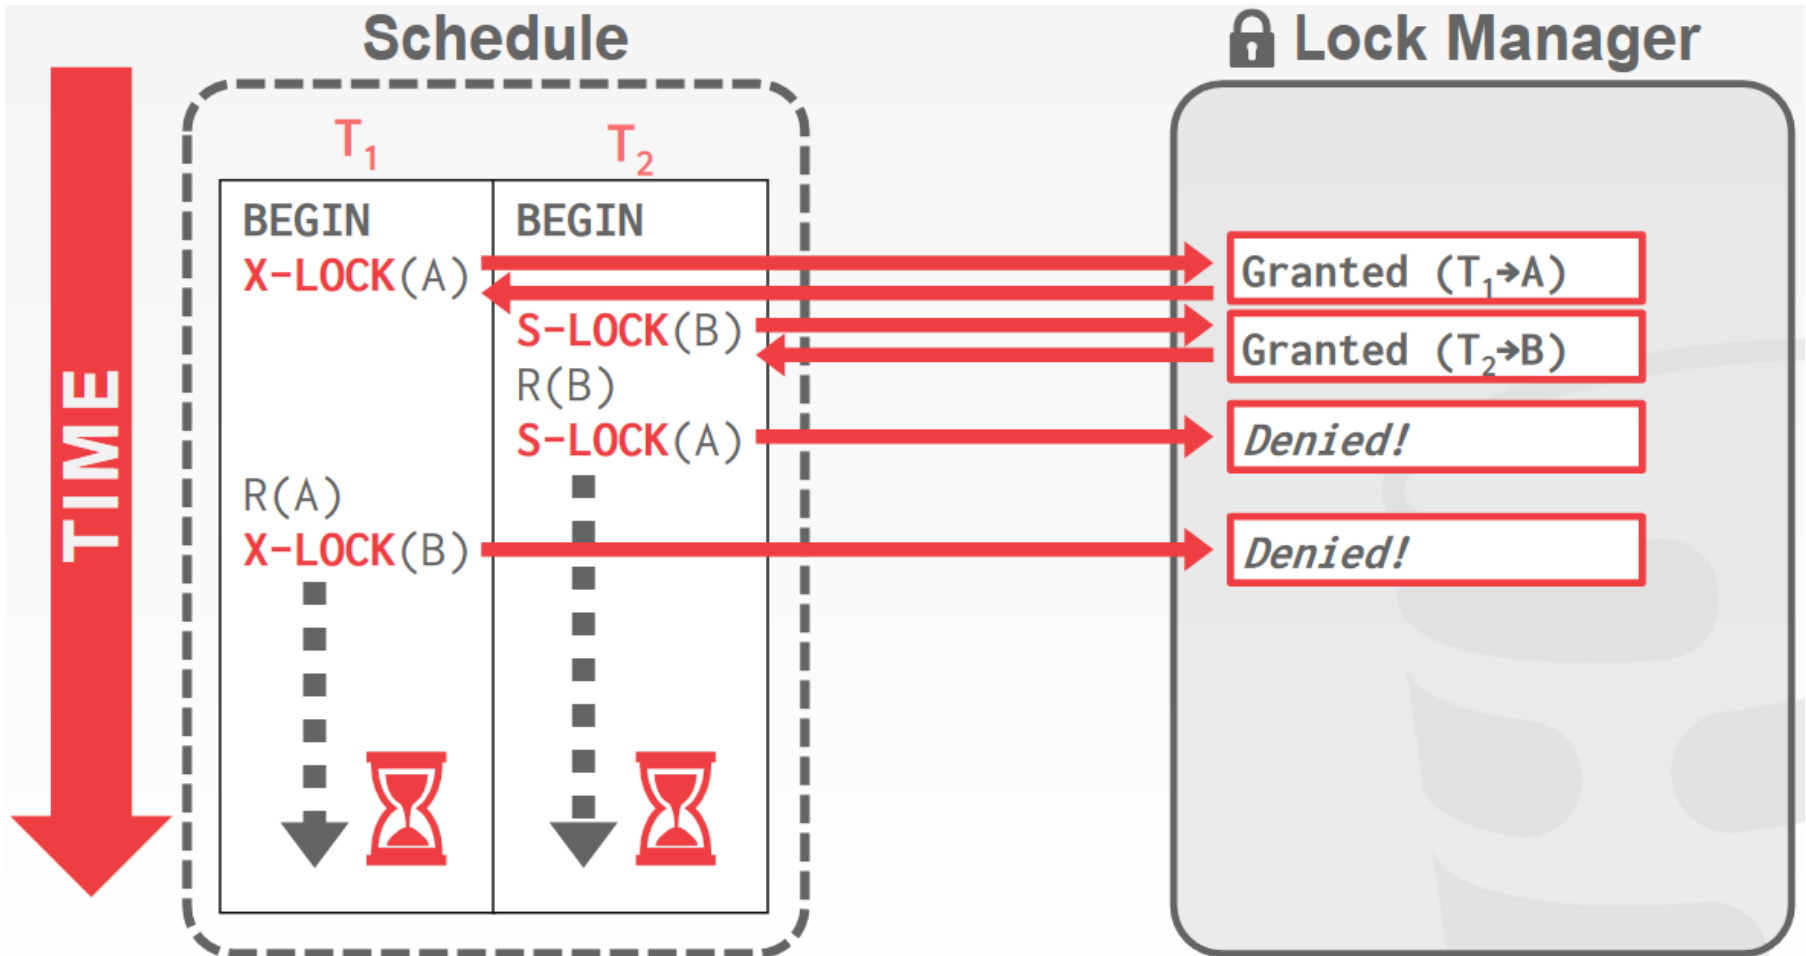


Deadlocks






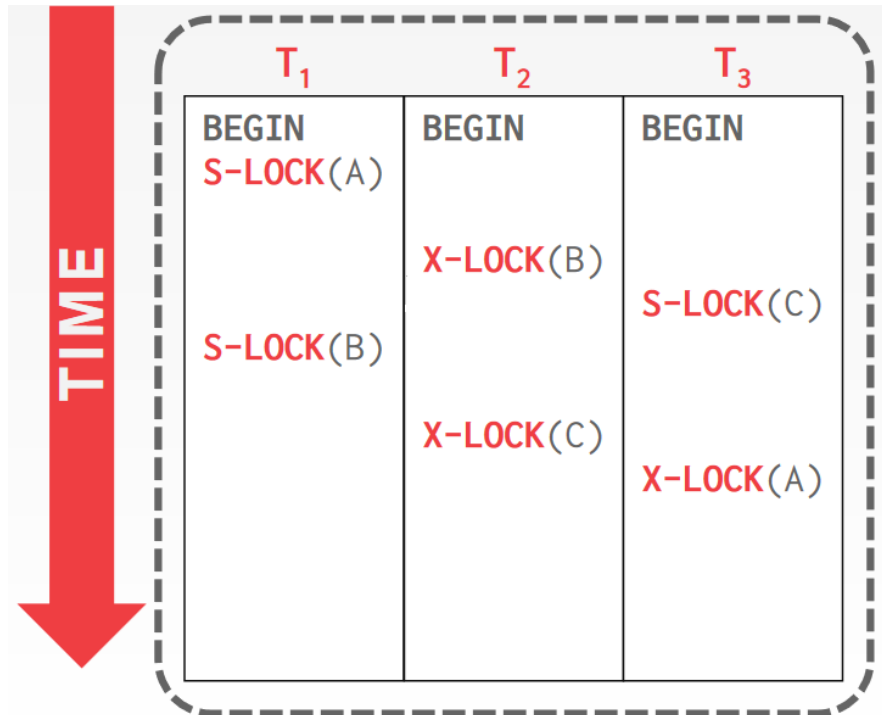
Deadlocks





Deadlock Detection

- We create a waits-for graph:
 - Each node is a transaction
 - Edge from T_i to T_j if T_i is waiting for T_j to release a lock
- Cycle in the graph means there's a deadlock





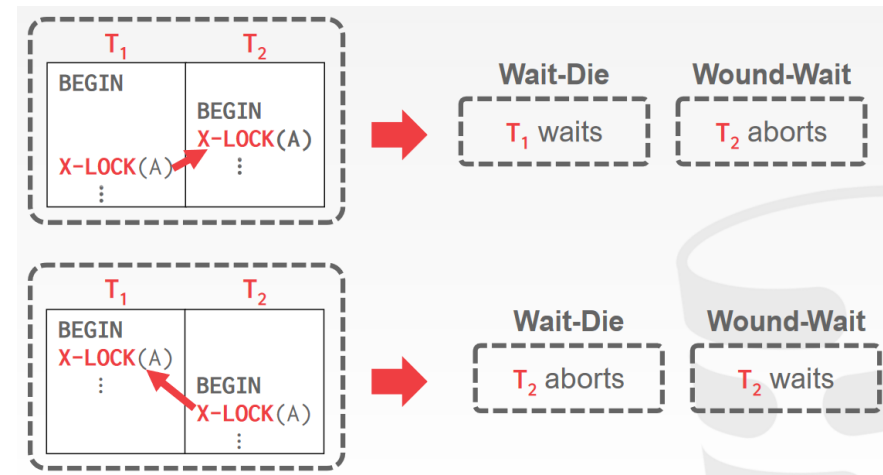
Deadlock Detection

- When the DBMS detects a deadlock using the waits-for graph, it will select a **victim transaction** to break the cycle.
 - The victim is forced to restart or abort.
- How to select the proper victim? Possible heuristics:
 - By age (e.g., lowest timestamp)
 - By progress (e.g., fewest queries executed)
 - # of objects locked
 - ...
- We should keep in mind whether a transaction has been selected as a victim before, to avoid **starvation**.



Deadlock Prevention

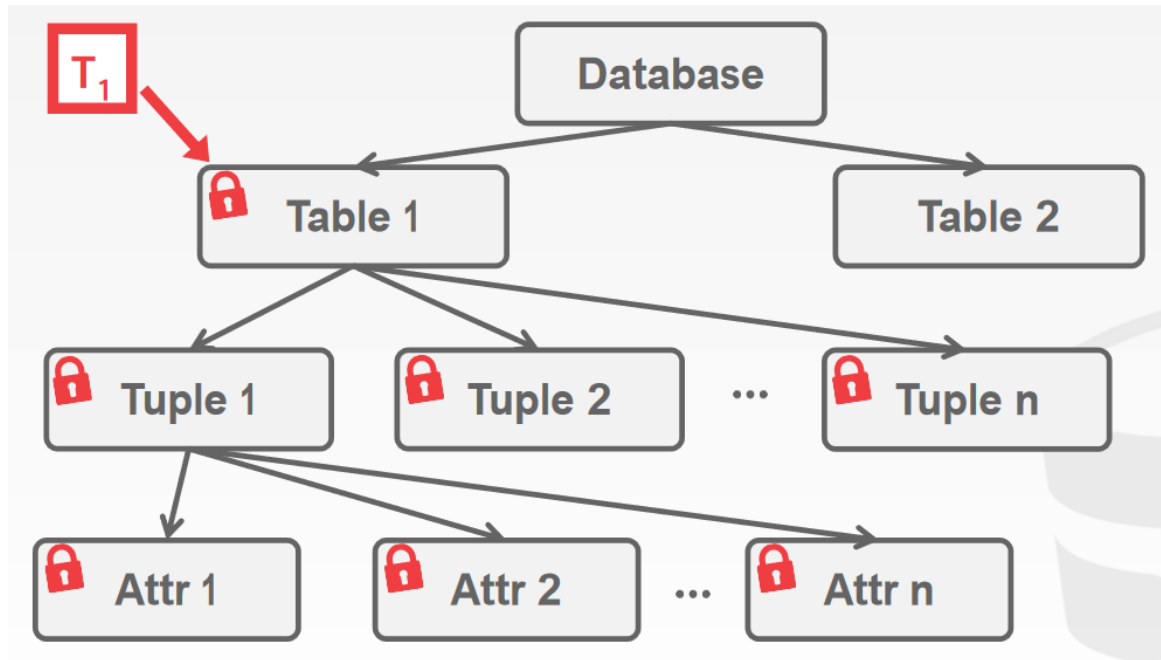
- When T_i tries to acquire a lock that is held by T_j , the DBMS "stops" one of them to prevent a deadlock.
 - Doesn't require a waits-for graph
- How to determine which transaction to stop?
 - Wait-Die**: If T_i started earlier than T_j , then T_i waits for T_j . Otherwise T_i aborts.
 - Wound-Wait**: If T_i started earlier than T_j , then T_j aborts and releases lock. Otherwise T_i waits.





Multi-Granularity Locking

- What are these **database objects** we have been locking?
 - Tables? Tuples? Attributes?
 - Locking too coarse: many conflicts, deadlocks
 - Locking too granular: can be inefficient






Multi-Granularity Locking

- An **intention lock** allows a higher level node to be locked without having to check all descendant nodes.
 - If a node is intention locked, then explicit locking is being done at a lower level in the tree.
- Five types of locks:
 - **S** lock: shared lock (same as before)
 - **X** lock: exclusive lock (same as before)
 - **IS** lock: intention-shared lock
 - You get IS at the parent, then S at one or more descendants
 - **IX** lock: intention-exclusive lock
 - You get IX at the parent, then X or S at one or more descendants
 - **SIX** lock: S + IX together



Example


Read Andy's record in **R**.





Example


Read Andy's record in **R**.





Example


Update Matt's record in R.





Example

Update Matt's record in R.





Compatibility


		T_2 Wants				
		IS	IX	S	SIX	X
T_1 Holds	IS	✓	✓	✓	✓	×
	IX	✓	✓	×	×	×
	S	✓	×	✓	×	×
	SIX	✓	×	×	×	×
	X	×	×	×	×	×

- If transactions' locks are not compatible, they cannot execute concurrently!



Example


Scan **R** and update a few tuples. **T₁**





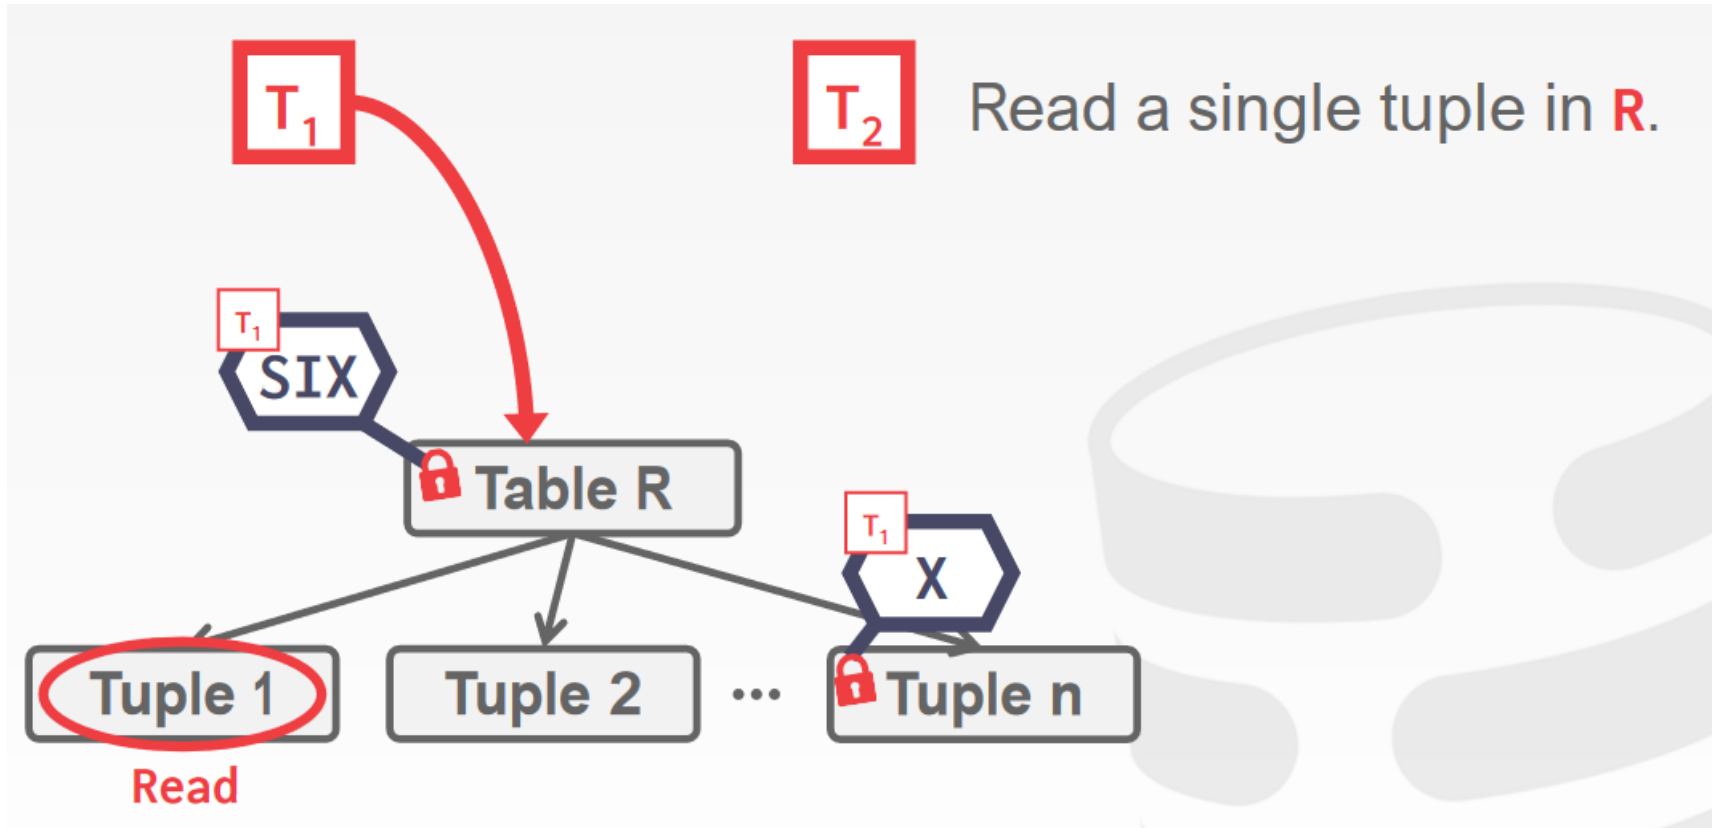
Example

Scan **R** and update a few tuples.



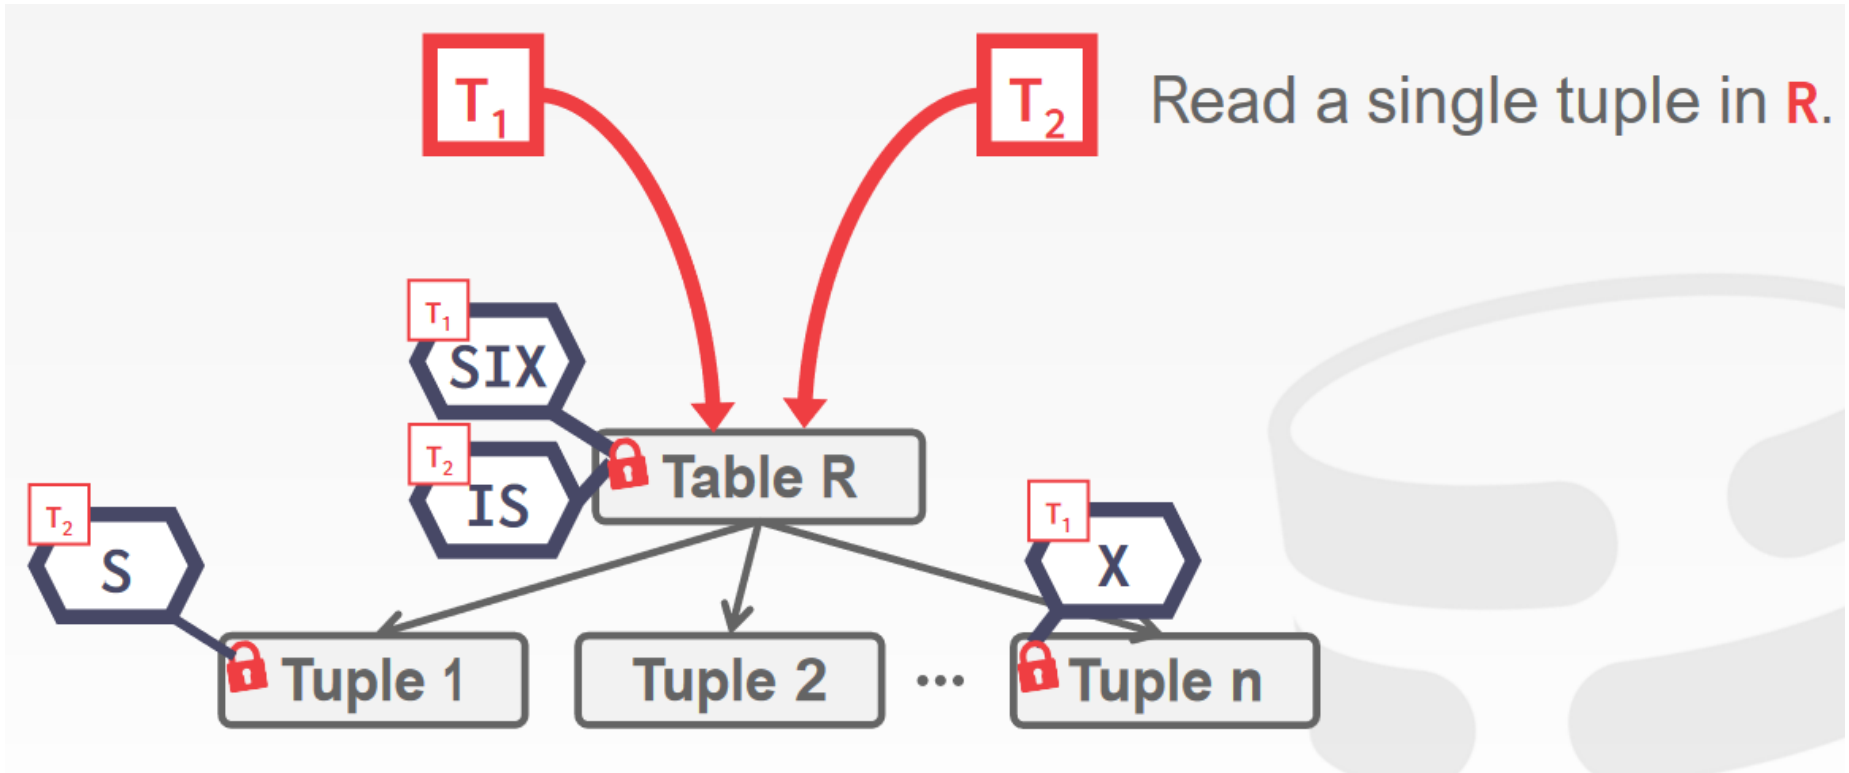


Example






Example






Example





Example

Scan all tuples in **R**.



Are the locks of these transactions compatible?



Locking in Practice

- Different DBMSs use different locking methods
- The DBMS can automatically take care of locking, unlocking, etc.
- But there **are** methods for explicit, manual locking
 - Depends on which DBMS you are using
 - **PostgreSQL**: Locking in SHARE, EXCLUSIVE, and other modes
 - **MySQL**: Locking in READ,WRITE modes

 PostgreSQL

ORACLE

IBM DB2

 Microsoft
SQL Server

 **MySQL**

```
LOCK TABLE <table> IN <mode> MODE;
```

```
SELECT 1 FROM <table> WITH (TABLOCK, <mode>);
```

```
LOCK TABLE <table> <mode>;
```